



PennState

# CMPSC 311 - Introduction to Systems Programming

Network Programming

Professors:  
Suman Saha

Slides are mostly by *Professor Patrick McDaniel*  
and *Professor Abutalib Aghayev*

If you ask if I got a UDP  
joke



is  
it still a UDP joke?


# What is a network?

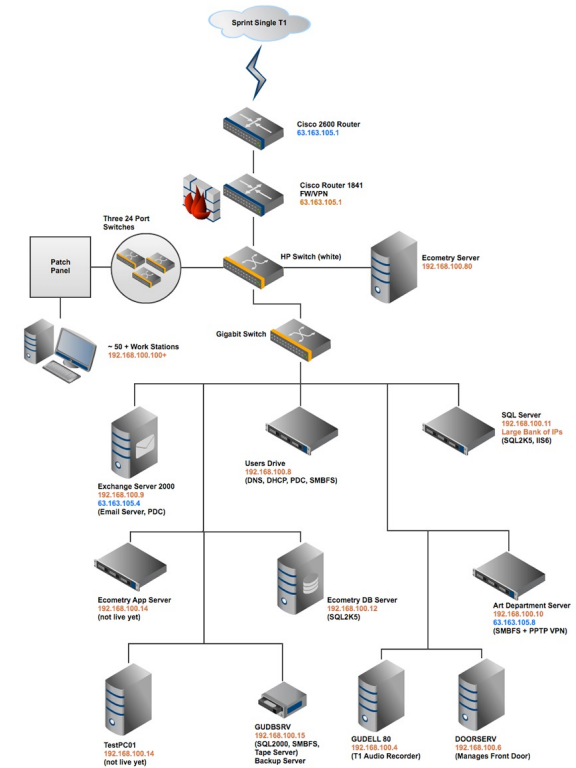


PennState

- A network is a collection of computing devices that share a transmission media

- Traditional wired networks (ethernet)
- High-speed backbone (fibre)
- Wireless (radio)
- Microwave
- Infrared


- The way the network is organized is called the network topology 



# The Internet



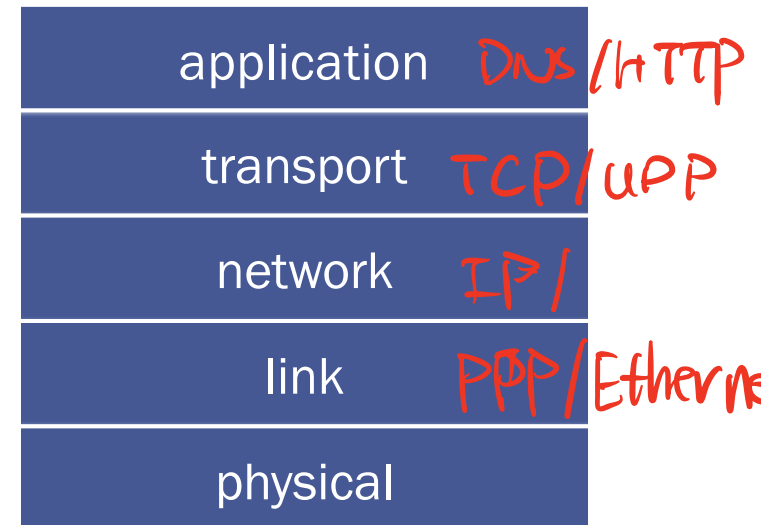
The Internet is an interconnected collection of many networks.



# Internet protocol stack (TCP/IP)



- **application**: supporting network applications
  - FTP, SMTP, HTTP, DNS
- **transport**: process-process data transfer
  - TCP, UDP
- **network**: routing of datagrams from source to destination
  - IP, routing protocols
- **link**: data transfer between neighboring network elements
  - PPP, Ethernet
- **physical**: bits “on the wire”



# Network vs. Web



PennState

- The network is a service ...
  - A conduit for data to be passed between systems.
  - Layers services (generally) to allow flexibility.
  - Highly scalable.
  - This is a public channel.
- The Web is an application
  - This is an application for viewing/manipulating content.
  - This can either be public (as in CNN's website), or private (as in enterprise internal HR websites).

# Networks Systems

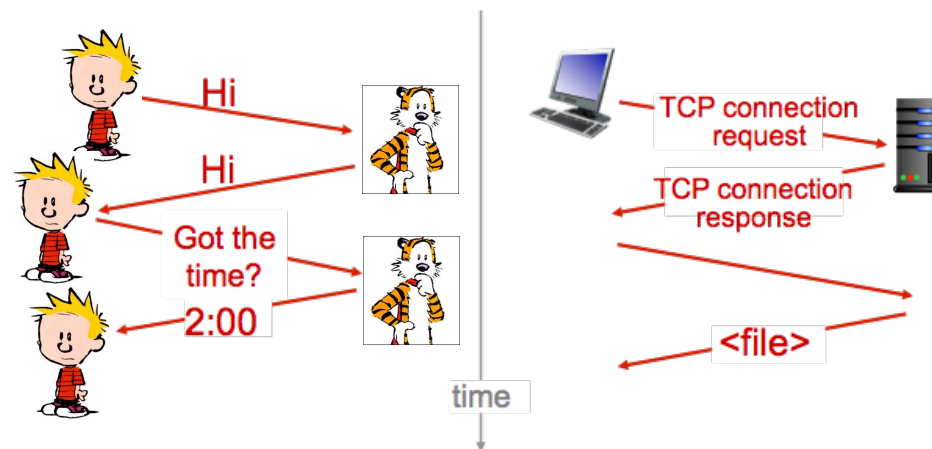


- Conceptually, think about network programming as two or more programs on the same or different computers talking to each other
  - The send messages back and forth
  - The “flow” of messages and the meaning of the message content is called the network protocol or just protocol



# What's a Protocol?

- Example: A human protocol and a computer protocol:



- Question: What are some other human protocols?

# Socket Programming

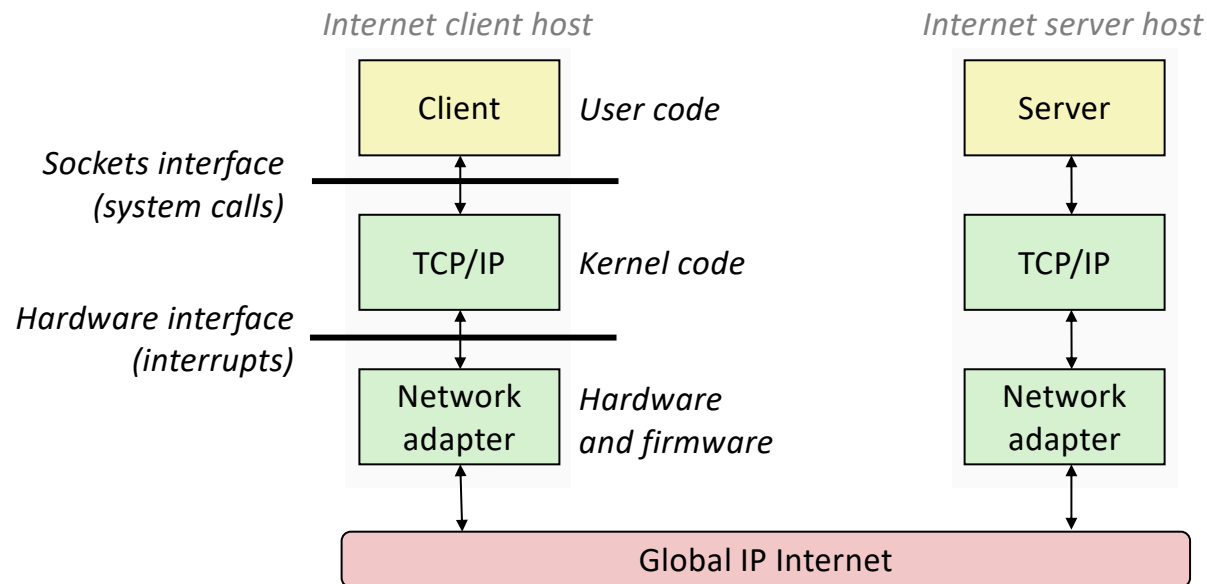


- Almost all meaningful careers in programming involve at least some level of network programming.
- Most of them involve [sockets](#) programming
  - Berkeley sockets originated in 4.2 BSD Unix circa 1983
    - it is the standard API for network programming
    - available on most OSs
  - POSIX socket API
    - a slight updating of the Berkeley sockets API
    - a few functions were deprecated or replaced
    - better support for multi-threading was added





# Hardware and Software Organization of Internet Application



- Bryant and O'Hallaron, Computer Systems: A Programmer's Approach

# IP addresses



- Every device on the Internet needs to have an address
  - Needs to be unique to make sure that it can be reached
- For IPv4, an IP address is a 4-byte tuple
  - e.g., 128.95.4.1 (80:5f:04:01 in hex)
- For IPv6, an IP address is a 16-byte tuple
  - e.g., 2d01:0db8:f188:0000:0000:0000:0000:1f33
  - 2d01:0db8:f188::1f33 in shorthand

# A Programmer's View of the Internet



- Hosts are mapped to a set of 32-bit **IP Address**
  - **146.186.145.12**
- The set of IP addresses mapped to a set of identifiers called Internet domain names
  - 146.186.145.12 is mapped to [www.eecs.psu.edu](http://www.eecs.psu.edu)
- A process on one Internet host can communicate with a process on another Internet host over a **connection**



# Recall file descriptors



- Remember open, read, write, and close?
  - POSIX system calls interacting with files
  - recall `open ( )` returns a `file descriptor`
    - an integer that represents an open file
    - inside the OS, it's an index into a table that keeps track of any state associated with your interactions, such as the file position
    - you pass the file descriptor into read, write, and close




# Networks and sockets



- UNIX makes all I/O look like file I/O
  - the good news is that you can use `read()` and `write()` to interact with remote computers over a network!
  - just like with files....
    - A program can have multiple network channels open at once
    - you need to pass `read()` and `write()` a file descriptor to let the OS know which network channel you want to write to or read from
  - The file descriptor used for network communications is a socket



# Pictorially




file descriptor	type	connected to?
0	pipe	stdin (console)
1	pipe	stdout (console)
2	pipe	stderr (console)
3	TCP socket	local: 128.95.4.33:80 remote: 44.1.19.32:7113
5	file	index.html
8	file	pic.png
9	TCP socket	local: 128.95.4.33:80 remote: 10.12.3.4:5544

OS's descriptor table

# Types of sockets



- Stream sockets
  - for connection-oriented, point-to-point, reliable bytestreams
    - uses **TCP**, SCTP, or other stream transports
- Datagram sockets
  - for connection-less, one-to-many, unreliable packets
    - uses **UDP** or other packet transports
- Raw sockets
  - for layer-3 communication (raw IP packet manipulation)



oriented

connection-oriented  
point-to-point  
reliable byte stream

# Stream (TCP) sockets




PennState

- Typically used for client / server communications
  - but also for other architectures, like peer-to-peer
- Client
  - an application that establishes a connection to a server
- Server
  - an application that receives connections from clients



1. establish connection



2. communicate



3. close connection



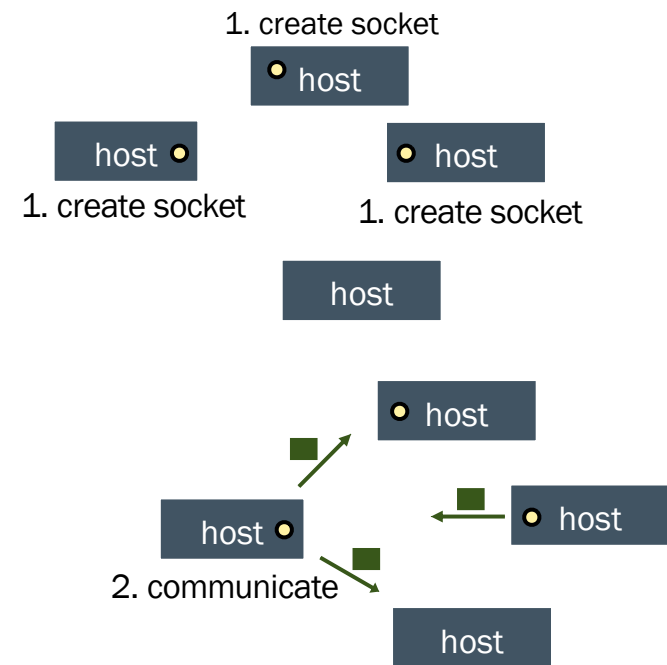
# Datagram (UDP) sockets



PennState

connection-less  
one-to-many  
unreliable byte streams

- Used less frequently than stream sockets
  - they provide no flow control, ordering, or reliability
  - really, provides best effort communication
- Often used as a building block
  - streaming media applications
  - sometimes, DNS lookups



**Note:** this is also called “*connectionless*” communication

# TCP connections



- Clients and servers communicate by sending streams of bytes over *connections*. Each connection is:
  - point-to-point: connects a pair of processes
  - full-duplex: data can flow in both directions at the same time
  - reliable: data send/receive order is preserved
- A *socket* is an endpoint of a connection
  - socket address is: `IPAddress:port pair`
- A *port* is a `16-bit integer` that identifies a process:
  - *ephemeral port*: assigned automatically by the client kernel when client makes a connection
  - *well-known port*: associated with some service provided by a server:  
(e.g. 80 is HTTP/Web)

短  
暂

# Network Ports




- Every computer has a numbered set of locations that represent the available “services” can be reached at
  - Ports are numbered 0 to 65535 → 16 bits
    - 0 to 1023 are called “well known” or “reserved” ports, where you need special (root) privileges to receive on these ports
  - Each transport (UDP/TCP) has its own list of ports
- Interesting port numbers
  - 20/21 - file transfer protocol (file passing)
  - 22 - secure shell (remote access)
  - 25 - Simple mail transfer protocol (email)
  - 53 - domain name service (internet naming)
  - 80 - HTTP (web)

# Anatomy of a connection



- A connection is uniquely identified by the socket addresses of its endpoints (*socket pair*)
  - (client IP:client port, server IP:server port)



51213 is an ephemeral port  
allocated by the kernel

80 is a well-known port  
associated with Web servers

# Programming a client



- We'll start by looking at the API from the point of view of a client connecting to a server over TCP
  - there are five steps:
    - figure out the address/port to connect to
    - create a socket
    - connect the socket to the remote server
    - `read()` and `write()` data using the socket
    - close the socket

1) get address  
2) Create socket  
3) connection  
4) read/write  
5) close



# inet\_aton()

1) Get Address



PennState

- The `inet_aton()` converts a IPv4 address into the UNIX structure used for processing:

```
int inet_aton(const char *addr, struct in_addr *inp);
```

- Where,
  - `addr` is a string containing the address to use (e.g., “166.84.7.99”)
  - `inp` is a pointer to the structure containing the UNIX internal representation of an address, used in later network communication calls

An example:

1. IPV4 to Binary: 166.84.7.99 → 10100110 01010100 00000111 01100011  
(2790524771)
2. Little Endian → Big Endian: 01100011 00000111 01010100 10100110
3. Binary → Decimal: 1661424806

`inet_aton()` returns 0 if failure!

# Putting it to use ...

1) Get Address



PennState

```
#include <stdlib.h>
#include <arpa/inet.h>

int main(int argc, char **argv) {
    struct sockaddr_in v4, sa; // IPv4
    struct sockaddr_in6 sa6; // IPv6

    // IPv4 string to sockaddr_in. (both ways)
    inet_aton("192.0.1.1", &(v4.sin_addr));
    inet_pton(AF_INET, "192.0.2.1", &(sa.sin_addr));
    // IPv6 string to sockaddr_in6.
    inet_pton(AF_INET6, "2001:db8:63b3:1::3490", &(sa6.sin6_addr));

    return( 0 );
}
```

Handwritten annotations: A red box highlights the two IPv4-related lines. A blue arrow points from the text "IPv4" to the boxed lines. Another blue arrow points from the text "IPv6" to the line containing the IPv6 address string.

# Getting back to strings?

1) Get Address



PennState

- The `inet_ntoa()` converts a UNIX structure for an IPv4 address into an ASCII string:

```
char *inet_ntoa(struct in_addr in);
```

```
struct sockaddr_in caddr, sa; // IPv4
struct sockaddr_in6 sa6; // IPv6
char astring[INET6_ADDRSTRLEN]; // IPv6

// Start by converting
inet_aton( "192.168.8.9", &caddr.sin_addr);
inet_pton(AF_INET, "192.0.2.1", &(sa.sin_addr));
inet_pton(AF_INET6, "2001:db8:63b3:1::3490", &(sa6.sin6_addr));

// Return to ASCII strings
inet_ntop(AF_INET6, &(sa6.sin6_addr), astring, INET6_ADDRSTRLEN);
printf( "IPv4 : %s\n", inet_ntoa(caddr.sin_addr) );
printf( "IPv6 : %s\n", astring );
```



# Domain Name Service (DNS)

1) Get Address



PennState

- People tend to use DNS names, not IP addresses
  - the sockets API lets you convert between the two
  - it's a complicated process, though:
    - a given DNS name can have many IP addresses
    - many different DNS names can map to the same IP address
      - an IP address will map onto at most one DNS names, and maybe none
    - a DNS lookup may require interacting with many DNS servers

**Note:** The “dig” Linux program is used to check DNS entries.

# Domain Name Service (DNS)

1) Get Address



PennState

- People

- the soc
- it's a co
- a
- m

- a

```
$ dig lion.cse.psu.edu

; <<>> DiG 9.9.2-P1 <<>> lion.cse.psu.edu
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 53447
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; MBZ: 0005 , udp: 4000
;; QUESTION SECTION:
;lion.cse.psu.edu.                IN      A

;; ANSWER SECTION:
lion.cse.psu.edu.                5       IN      A      130.203.22.184

;; Query time: 38 msec
;; SERVER: 127.0.1.1#53(127.0.1.1)
;; WHEN: Tue Nov 12 14:02:11 2013
;; MSG SIZE rcvd: 61
```

**Note:** The “dig” Linux program is used to check DNS entries.

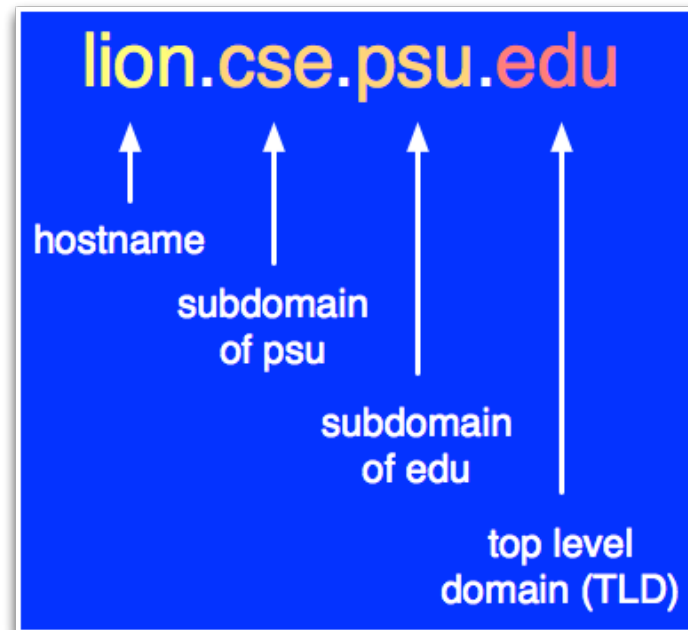
# The FQDN

1) Get Address



PennState

- Every system that is supported by DNS has a unique **fully qualified domain name** (FQDN)



# DNS hierarchy

1) Get Address



PennState



# Resolving DNS names

1) Get Address



PennState

- The POSIX way is to use `getaddrinfo()` — *complicated way*
- a pretty complicated system call; the basic idea...
  - set up a “hints” structure with constraints you want respected
    - e.g., IPv6, IPv4, or either
  - indicate which host and port you want resolved
    - host: a string representation; DNS name or IP address
  - returns a list of results packet in an “`addrinfo`” struct
  - free the `addrinfo` structure using `freeaddrinfo()`

# DNS resolution (the easy way)



PennState

1) Get Address

- The `gethostbyname()` uses DNS to look up a name and return the host information

```
struct hostent *gethostbyname(const char *name);
```

- Where,

→ DNS

- ▶ `name` is a string containing the host/domain name to find
- ▶ `hostent` is a structure with the given host name. Here name is either a hostname, or an IPv4 address in standard dot notation. The structure includes:
  - ▶ `hostent->h_name` (fully qualified domain name ~~name~~) FQDN
  - ▶ `hostent->h_addr_list` (list of pointers to IP addresses)

# DNS resolution (the easy way)

1) Get Address



PennState

- The `gethostbyname()` uses DNS to look up a name and return the host information

```
struct hostent *gethostbyname(const char *name);
```

- Where,

- ▶ `name` is

- ▶ `hostent`

- or an IPv

- ▶ `host`

- ▶ `host`

```
char *hn = "lion.cse.psu.edu";
struct hostent * hstinfo;
struct in_addr **addr_list;
if ((hstinfo = gethostbyname(hn)) == NULL) {
    return -1;
}

addr_list = (struct in_addr **)hstinfo->h_addr_list;
printf( "DNS lookup [%s] address [%s]\n", hstinfo->h_name,
        inet_ntoa(*addr_list[0]) );
```


```
$ ./network
DNS lookup [lion.cse.psu.edu] address [130.203.22.184]
$
```

hostname,

# Programming a client



- We'll start by looking at the API from the point of view of a client connecting to a server over TCP
  - there are five steps:
    - figure out the address/port to connect to
    - create a socket
    - connect the socket to the remote server
    - read() and write() data using the socket
    - close the socket





# Creating a socket

2) Create the socket



PennState

- The `socket()` function creates a file handle for use in communication:

```
int socket(int domain, int type, int protocol);
```

- Where,

- ▶ `domain` is the communication domain
  - ▶ `AF_INET` (IPv4), `AF_INET6` (IPv6)
- ▶ `type` is the communication semantics (stream vs. datagram)
  - ▶ `SOCK_STREAM` is stream (using TCP by default)
  - ▶ `SOCK_DGRAM` is datagram (using UDP by default)
- ▶ `protocol` selects a protocol from available (not used often)

**Note:** creating a socket doesn't connect to anything

# Creating a socket

2) Create the socket



PennState

- The `socket()` function creates a file handle for use in communication:

```
int socket(int domain, int type, int protocol);
```

- Where,

- ▶ `domain`

- ▶ `AF_INET`

- ▶ `type`

- ▶ `SOCK_STREAM`

- ▶ `SOCK_DGRAM` is datagram (using UDP by default)

- ▶ `protocol` selects a protocol from available (not used often)

```
// Create the socket
int sockfd;
sockfd = socket(AF_INET, SOCK_STREAM, 0);
if (sockfd == -1) {
    printf( "Error on socket creation [%s]\n", strerror(errno) );
    return( -1 );
}
```

**Note:** creating a socket doesn't connect to anything

# Specifying an address

3) Connect to server



PennState

- The next step is to create an address to connect to by specifying the address and port in the proper form.
  - protocol family (`addr.sin_family`)
  - port (`addr.sin_port`)
  - IP address (`addr.sin_addr`)

```
// Variables
char *ip = "127.0.0.1";
unsigned short port = 16453;
struct sockaddr_in caddr;

// Setup the address information
caddr.sin_family = AF_INET;
caddr.sin_port = htons(port);
if ( inet_aton(ip, &caddr.sin_addr) == 0 ) {
    return( -1 );
}
```

# Network byte order

3) Connect to server




PennState

- When sending data over a network you need to convert your integers to be in network byte order, and back to host byte order upon receive:

```
uint64_t htonll64(uint64_t hostlong64);  
uint32_t htonl(uint32_t hostlong);  
uint16_t htons(uint16_t hostshort);  
uint64_t ntohll64(uint64_t hostlong64);  
uint32_t ntohl(uint32_t netlong);  
uint16_t ntohs(uint16_t netshort);
```

- Where each of these functions receives a NBO or HBO 64/ 32 /16 byte and converts it to the other.

Network byte order (Big Endian)



# connect()

3) Connect to server



PennState

- The connect() system call connects the socket file descriptor to the specified address

```
int connect(int sockfd, const struct sockaddr *addr,  
            socklen_t addrlen);
```

- Where,
  - sockfd is the socket (file handle) obtained previously
  - addr - is the address structure
  - addlen - is the size of the address structure
  - Returns 0 if successfully connected, -1 if not

```
if ( connect(sockfd, (const struct sockaddr *)&caddr,  
             sizeof(caddr)) == -1 ) {  
    return( -1 );  
}
```

# Reading and Writing

4) Send and  
receive data



PennState

- Primitive reading and writing only process only blocks of opaque data:

```
ssize_t write(int fd, const void *buf, size_t count);  
ssize_t read(int fd, void *buf, size_t count);
```

- Where `fd` is the file descriptor, `buf` is an array of bytes to write from or read into, and `count` is the number of bytes to read or write
- The value returned is the number of bytes read or written.
  - Be sure to always check the result
- On reads, you are responsible for supplying a buffer that is large enough to put the output into.
  - look out for memory corruption when buffer is too small ...

# close()

5) Close the  
socket



PennState

- `close()` closes the connection and deletes the associated entry in the operating system's internal structures


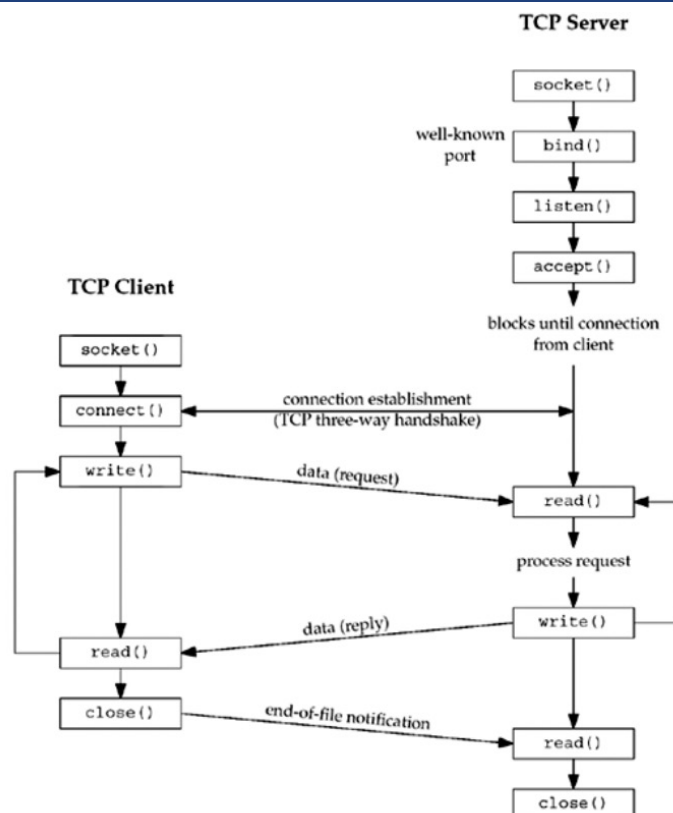
```
close( socket_fd );  
socket_fd = -1;
```

**Note:** set handles to `-1` to avoid use after close.

# Elementary TCP client/server



PennState





# Programming a server

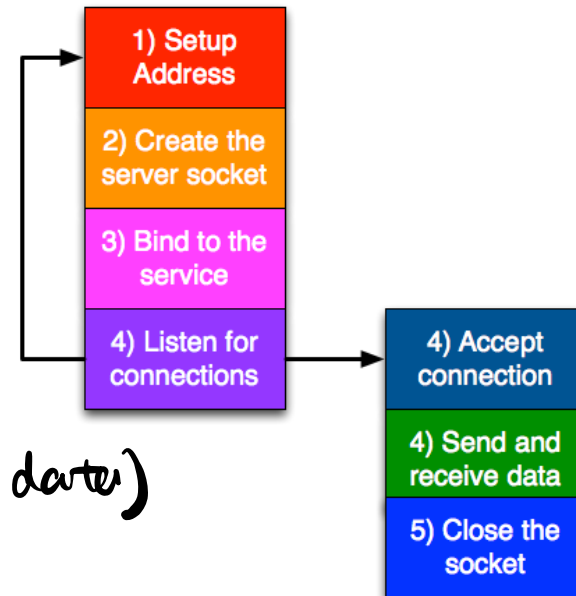


- Now we'll look at the API from the point of view of a server who receives connections from clients

- there are seven steps:

- figure out the port to "bind" to
- create a socket
- bind the service
- begin listening for connections
- receive connection
- read() and write() data (send or receive data)
- close the socket

*server connection*



# Setting up a server address

1) Setup  
Address



PennState

- All you need to do is specify the service port you will use for the connection:

```
struct sockaddr_in saddr;  
saddr.sin_family = AF_INET;  
saddr.sin_port = htons(16453);  
saddr.sin_addr.s_addr = htonl(INADDR_ANY);
```

- However, you don't specify an IP address because you are receiving connections at the local host.
  - Instead you use the special "any" address

`htonl(INADDR_ANY)`

**Next:** creating the socket is as with the client

2) Create the  
server socket

*Handwritten:*  
`int socket(domain, type, protocol)`  
          ↓                  ↓  
      IPv4/IPv6      sock.STREAM

# Binding the service

3) Bind to the service



PennState

- The `bind()` system call associates a socket with the server connection (e.g., taking control of HTTP)

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t*addrlen)
```

- Where,
  - `sockfd` is the socket (file handle) obtained previously
  - `addr` - is the address structure
  - `addrlen` - is the size of the address structure
  - Returns 0 if successfully connected, -1 if not

↓  
similar to connection() in client side.

```
if ( bind(sock, (const struct sockaddr *)&saddr,  
          sizeof(saddr)) == -1 ) {  
    return( -1 );  
}
```

# Listening for connections

4) Listen for connections



PennState

- The `listen()` system call tells the OS to receive connections for the process

```
int listen(int sockfd, int backlog);
```

- Where,

- `sockfd` is the socket (file handle) obtained previously
- `backlog` - is the number of connections to queue
  - A program may process connections as fast as it wants, and the OS will hold the client in a waiting state until you are ready
  - Beware of waiting too long (timeout)

```
if ( listen(sock, 5) == -1 ) {  
    return( -1 );  
}
```

# Accepting connections

4) Accept connection



PennState

- The `accept ( )` system call receives the connection from the client:  
`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`
- Where,
  - `sockfd` is the socket (file handle) obtained previously
  - `addr` - is the address structure for client (filled in)
  - `addrlen` - is the size of the address structure
  - Returns the new socket handle or -1 if failure

`bind ( )`  
`listen ( )`  
`accept ( )`

# Accepting connections

4) Accept connection



PennState

- The `accept ( )` system call receives the connection from the client:  
`int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);`
- Where,
  - `sockfd` is the socket (file handle) obtained previously
  - `addr` - is the address structure for client (filled in)
  - `addrlen` - is the size of the address structure
- Returns the new socket handle or -1 if failure

```
inet_len = sizeof(caddr);
if ( (client = accept(server, (struct sockaddr *)&caddr, inet_len )) == -1 )
{
    printf( "Error on client accept [%s]\n", strerror(errno) );
    close(server);
    return( -1 );
}
printf( "Server new client connection [%s/%d]",
        inet_ntoa(caddr.sin_addr), caddr.sin_port );
```

# The rest ...



- From the server perspective, receiving and sending on the newly received socket is the same as if it were a client
  - `read()` and `write()` for sending and receiving
  - `close()` for closing the socket

4) Send and  
receive data

5) Close the  
socket

# Putting it all together (client)



PennState

```
int client_operation( void ) {  
  
    int socket_fd;  
    uint32_t value;  
    struct sockaddr_in caddr;  
    char *ip = "127.0.0.1";  
  
    caddr.sin_family = AF_INET;  
    caddr.sin_port = htons(16453);  
    if ( inet_aton(ip, &caddr.sin_addr) == 0 ) {  
        return( -1 );  
    }  
  
    socket_fd = socket(PF_INET, SOCK_STREAM, 0);  
    if (socket_fd == -1) {  
        printf( "Error on socket creation [%s]\n", strerror(errno) );  
        return( -1 );  
    }  
  
    if ( connect(socket_fd, (const struct sockaddr *)&caddr, sizeof(caddr)) == -1 ) {  
        printf( "Error on socket connect [%s]\n", strerror(errno) );  
        return( -1 );  
    }  
  
    value = htonl( 1 );  
    if ( write( socket_fd, &value, sizeof(value)) != sizeof(value) ) {  
        printf( "Error writing network data [%s]\n", strerror(errno) );  
        return( -1 );  
    }  
    printf( "Sent a value of [%d]\n", ntohl(value) );  
  
    if ( read( socket_fd, &value, sizeof(value)) != sizeof(value) ) {  
        printf( "Error reading network data [%s]\n", strerror(errno) );  
        return( -1 );  
    }  
    value = ntohl(value);  
    printf( "Receivd a value of [%d]\n", value );  
  
    close(socket_fd); // Close the socket  
    return( 0 );  
}
```



# Putting it all together (server)



PennState

```
int server_operation( void ) {

    int server, client;
    uint32_t value, inet_len;
    struct sockaddr_in saddr, caddr;

    saddr.sin_family = AF_INET;
    saddr.sin_port = htons(16453);
    saddr.sin_addr.s_addr = htonl(INADDR_ANY);

    server = socket(PF_INET, SOCK_STREAM, 0);
    if (server == -1) {
        printf( "Error on socket creation [%s]\n", strerror(errno) );
        return( -1 );
    }

    if ( bind(server, (struct sockaddr *)&saddr, sizeof(saddr)) == -1 ) {
        printf( "Error on socket bind [%s]\n", strerror(errno) );
        return( -1 );
    }

    if ( listen( server, 5 ) == -1 ) {
        printf( "Error on socket listen [%s]\n", strerror(errno) );
        return( -1 );
    }
}
```

## ... Together (server, part 2)



PennState

```
while ( 1 ) {
    inet_len = sizeof(caddr);
    if ( (client = accept( server, (struct sockaddr *)&caddr, &inet_len )) == -1 ) {
        printf( "Error on client accept [%s]\n", strerror(errno) );
        close(server);
        return( -1 );
    }
    printf( "Server new client connection [%s/%d]", inet_ntoa(caddr.sin_addr), caddr.sin_port );
    if ( read( client, &value, sizeof(value)) != sizeof(value) ) {
        printf( "Error writing network data [%s]\n", strerror(errno) );
        close(server);
        return( -1 );
    }
    value = ntohl(value);
    printf( "Receivd a value of [%d]\n", value );
    value++;
    value = htonl(value);
    if ( write( client, &value, sizeof(value)) != sizeof(value) ) {
        printf( "Error writing network data [%s]\n", strerror(errno) );
        close(server);
        return( -1 );
    }
    printf( "Sent a value of [%d]\n", value );
    close(client); // Close the socket
}
return( 0 );
}
```