



PennState


CMPSC 311 - Introduction to Systems Programming

Signals

Professors:

Suman Saha

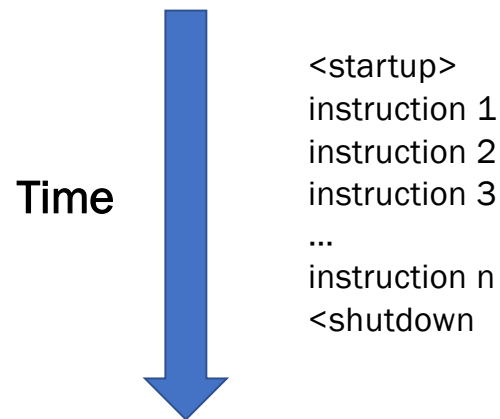
(Slides are mostly by *Professor Patrick McDaniel*
and *Professor Abutalib Aghayev*)



Control Flow

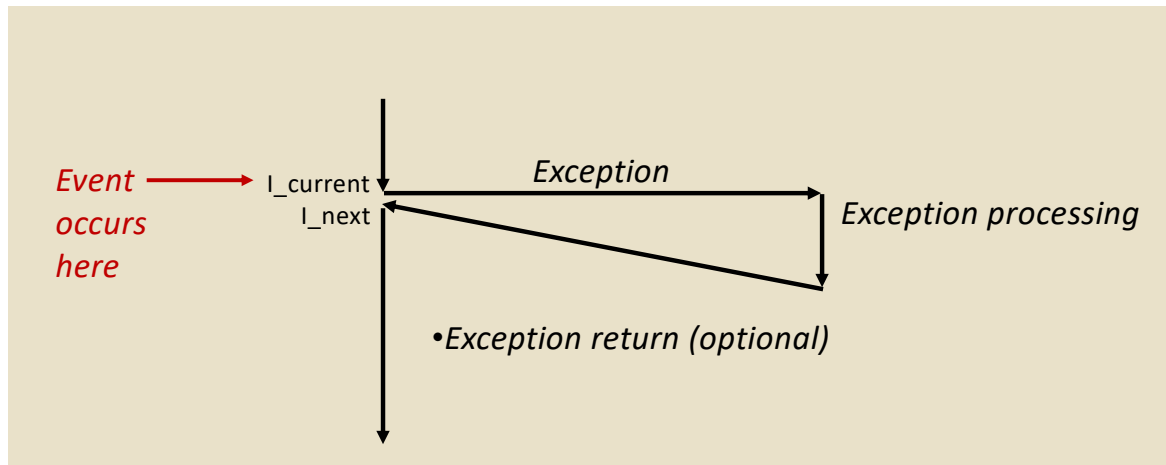


- Processors do only one thing:
 - from startup to shutdown, a CPU simply reads and executes a sequence of instructions, one at a time
 - This sequence is the CPU's control flow (or flow of control)



Exceptional Control Flow

- Exceptional control flow enables a system to react to an event



Exceptional Control Flow




- Mechanisms exist at all levels of a computer system for exceptional control
- Low-level mechanisms
 - **Exceptions**
 - Examples: interrupts, traps, faults, and aborts
 - Implemented using combination of hardware and OS software
- High-level mechanisms
 - **Process context switch** (implemented by OS software and hardware timer)
 - **Signals** (implemented by OS software)
 - **Nonlocal jumps**: `setjmp()` and `longjmp()` (implemented by C runtime library)

UNIX Signals



- A **signal** is a special message sent through the OS to tell a process (or thread) of some command or event
- The process execution stops and special “**signal handler**” code runs.
- The process can resume operation after the signal handling is complete.



Signal types (abbreviated)



```
/* Signals */
#define SIGHUP          1      /* Hangup (POSIX).  */
#define SIGINT          2      /* Interrupt (ANSI). */
#define SIGQUIT         3      /* Quit (POSIX).   */
#define SIGABRT         6      /* Abort (ANSI).   */
#define SIGFPE          8      /* Floating-point exception (ANSI). */
#define SIGKILL         9      /* Kill, unblockable (POSIX). */
#define SIGSEGV       11      /* Segmentation violation (ANSI). */
#define SIGTERM       15      /* Termination (ANSI). */
#define SIGSTKFLT     16      /* Stack fault.    */
#define SIGCHLD         17      /* Child status has changed (POSIX). */
#define SIGCONT         18      /* Continue (POSIX). */
#define SIGSYS          31      /* Bad system call. */
```

Signals as process control



- The operating system use signals to control process behavior

- Signals are sent on errors

```
#define SIGILL      4      /* Illegal instruction (ANSI).  */
#define SIGTRAP     5      /* Trace trap (POSIX).  */
#define SIGIOT      6      /* IOT trap (4.2 BSD).  */
#define SIGBUS      7      /* BUS error (4.2 BSD).  */
#define SIGFPE      8      /* Floating-point exception (ANSI).  */
#define SIGSEGV     11     /* Segmentation violation (ANSI).  */
```

- Signals can be used by other applications too

```
#define SIGUSR1     10     /* User-defined signal 1 (POSIX).  */
#define SIGUSR2     12     /* User-defined signal 2 (POSIX).  */
```

- Control the process execution

```
#define SIGKILL      9      /* Kill, unblockable (POSIX).  */
#define SIGCONT     18      /* Continue (POSIX).  */
#define SIGSTOP     19      /* Stop, unblockable (POSIX).  */
```

Process IDs



- Every process running on the OS is given a unique process ID (PID)
 - This is what is used in the OS and for process control to reference that specific running program instance.
- To find a process ID for a program, use the **ps** utility
 - The **ps** stands for “process status”

```
$ ps -U mcdaniel
  PID TTY          TIME CMD
 30908 ?            00:00:00 gnome-keyring-d
 30919 ?            00:00:00 gnome-session
 30964 ?            00:00:00 ssh-agent
 30967 ?            00:00:00 dbus-launch
 30968 ?            00:00:01 dbus-daemon
 30978 ?            00:00:00 at-spi-bus-laun
 30982 ?            00:00:00 dbus-daemon
 30985 ?            00:00:00 at-spi2-registr
 30999 ?            00:00:02 gnome-settings-
31009 ?            00:00:00 pulseaudio
31011 ?            00:00:00 gvfsd
31017 ?            00:00:00 gvfsd-fuse
31031 ?            00:02:43 compiz
31041 ?            00:00:00 dconf-service
31044 ?            00:00:00 gnome-fallback-
31045 ?            00:00:06 nautilus
31047 ?            00:00:01 nm-applet
31048 ?            00:00:41 vmtoolsd
31049 ?            00:00:00 polkit-gnome-au
31064 ?            00:00:00 gvfs-udisks2-vo
31079 ?            00:00:00 gvfs-gphoto2-vo
31083 ?            00:00:00 gvfs-afc-volume
31090 ?            00:00:00 gvfs-mtp-volume
...
```


kill



PennState

- Kill is a program than sends signals to processes.

`kill [-<sig>] <pid>`

- Where **<sig>** is the signal number and **<pid>** is the process ID of the running program you want to send the signal.
 - If no **SIGNUM** is given, then **SIGTERM** is used by default.

PID → *program*

```
$ ps -U mcdaniel
57613 pts/4    00:00:00 signals
$ kill -1 57613
$ kill -2 57613
$ kill -9 57613
```

<sig> ↓

1: SIGHUP
2: SIGINT

```
$ ./signals
Sleeping ...zzzzz ....
Signal handler got a SIGHUP!
Signals received : 1
Woken up!!
Sleeping ...zzzzz ....
Signal handler got a SIGINT!
Signals received : 2
Woken up!!
Sleeping ...zzzzz ....
Killed
```

9: SIGKILL

SIGTERM vs. SIGKILL



- **SIGTERM** interrupts the program and asks it to shut down, which it should.
 - Sometimes this does not work (for instance when the process is in a locked state)
 - It is often desirable to add a signal handler to handle the SIGTERM, so that it can gracefully shut down the process, cleanup memory, close files, etc.
- **SIGKILL** kills the process
 - Can lead to inconsistent state, because there is no opportunity to gracefully shutdown the process.

Definition: the term *graceful shutdown* refers to the proper and complete sync with secondary storage, disposal of resources, and normal termination.

send signals to all instances of a particular program

killall



PennState

- Killall is a program than sends signals to all instances of a particular program.

`killall [-<sig>] <name>`

- Where **<sig>** is the signal number and **<name>** is the name of running program you want to send the signal.
 - If no **SIGNAL** is given, then **SIGTERM** is used by default.

```
$ killall -1 signals
$ killall -2 signals
$ killall -SIGKILL signals
```

```
$ ./signals
Sleeping ...zzzzz ...
Signal handler got a SIGHUP!
Signals received : 1
Woken up!!
Sleeping ...zzzzz ...
Signal handler got a SIGNINT!
Signals received : 2
Woken up!!
Sleeping ...zzzzz ...
Killed
```

raise()



- raise allows a process to send signals to itself.

```
int raise(int sig);
```

- There are a range of reasons why a process might want to do this.
 - Suspend itself (SIGSTOP)
 - Kill itself (SIGKILL)
 - Reset its configuration (SIGHUP)
 - User defined signals (SIGUSR1..)

```
void suicide_signal(void) {  
    raise(SIGKILL);  
    return; // This will never be reached  
}
```



User-defined signal handlers



- You can create your own signal handlers simply by creating a function

```
void <fname>( int <var name> )
```

- and passing a function pointer to the function

```
sighandler_t signal(int signum, sighandler_t handler);
```

- Thereafter, whenever a signal of the type `signo` is raised, your program is called instead of the default handler.

```
void signal_handler(int no) {  
    printf("Sig handler got a [%d]\n", no);  
    return;  
}
```

```
signal(SIGHUP, signal_handler);  
signal(SIGINT, signal_handler);
```

Function pointers



- A function pointer is a pointer to a function that can be assigned, passed as parameters, and called

`<return> (*<var>) (<params>);`

- `<return>` is the return type of the function
- `<var>` is the variable names
- `<params >` are the parameters , separated by commas

```
int myfunc(int i) {  
    printf("Got into function with %d\n", i);  
    return 0;  
}  
  
int main( void ) {  
    int (*func) (int);  
    func = myfunc; //set variable name to function  
    func(7);  
    return 0;  
}
```

```
$ ./signals  
Got into function with 7  
$
```

An alternate approach



- The `sigaction()` system call changes the action taken by a process on receipt of a specific signal.

```
int sigaction(int signum, const struct sigaction *act, struct sigaction *oldact);
```

- Where:

- `signum` - is the signal number to be handled
- `act` - is a structure containing information about the `new handler`, NULL means ignore the signal
- `oldact` - is a pointer to the `previously assigned handler` as assigned in call to function

```
struct sigaction new_action, old_action;  
new_action.sa_handler = signal_handler;  
new_action.sa_flags = SA_NODEFER | SA_ONSTACK;  
sigaction(SIGINT, &new_action, &old_action);
```

```
signalhandler_t signal (int signum, signalhandler_t handler);  
int sigaction (int signum, const struct sigaction *new-act,  
              struct sigaction *old-act);
```

Why another API?



- Many argue that the `sigaction` function is better:
 - The `signal()` function does **not block** other signals from arriving while the current handler is executing; `sigaction()` can **block** other signals until the current handler returns.
 - The `signal()` function resets the signal action back to `SIG_DFL` (default) for almost all signals.
 - Better tuning of signals/controls of process through flags
 - `SA_NODEFER` - don't suspend signals while in handler
 - `SA_ONSTACK` - provide alternate stack for signal handler
 - `SA_RESETHAND` - Restore the signal action to the default upon entry to the signal handler.

Note: *In general, `sigaction` is preferred over `signal`.*

Putting it all together ...



PennState

```
void signal_handler(int no) {
    printf("Signal received : %d\n", no);
    if (no == SIGHUP) {
        printf("Signal handler got a SIGHUP!\n");
    } else if (no == SIGINT) {
        printf("Signal handler got a SIGNINT!\n");
    }
    return;
}

void cleanup_handler(int no) {
    printf("Killed");
    exit(0);
}

int main(void) {
    struct sigaction new_action, old_action; // Setup the signal actions
    new_action.sa_handler = signal_handler;
    new_action.sa_flags = SA_NODEFER | SA_ONSTACK;
    sigaction( SIGINT, &new_action, &old_action );
    signal( SIGHUP, signal_handler ); // Setup the signal handlers
    signal( SIGTERM, cleanup_handler );

    while (1) {
        printf( "Sleeping ...zzzzz ....\n" );
        select( 0, NULL, NULL, NULL, NULL );
        printf( "Woken up!!\n" );
    }

    // Return successfully
    return 0;
}
```

```
$ ./signals
Sleeping ...zzzzz ....
Signal received : 1
Signal handler got a SIGHUP!
Woken up!!
Sleeping ...zzzzz ....
Signal received : 2
Signal handler got a SIGNINT!
Woken up!!
Sleeping ...zzzzz ....
Killed
```