

Blocking Solutions

Where are we?

- Disable Interrupts
 - Effectively stops scheduling other activities.
- Busy-wait/spinlock Solutions
 - Pure software solutions *mux progress bounded wait requirement*
 - Integrated hardware-software solutions
- Blocking Solutions

Spinning vs. Blocking

- In the previous solns., we busy-waited for some condition to change.
- This change should be affected by some other activity.
- We are “presuming” that this other activity will eventually get the CPU (some kind of pre-emptive scheduler).
- This can be inefficient because:
 - You are wasting the rest of your time quantum in busy-waiting
 - Sometimes, your programs may not work! (if the OS scheduler is not pre-emptive).

- In blocking solutions, you relinquish the CPU at the time you cannot proceed, i.e., you are put in the blocked queue.
- It is the job of the activity changing the condition to wake you up (i.e., move you from blocked back to ready queue).
- This way you do not unnecessarily occupy CPU cycles.

Example Blocking Implementation

```

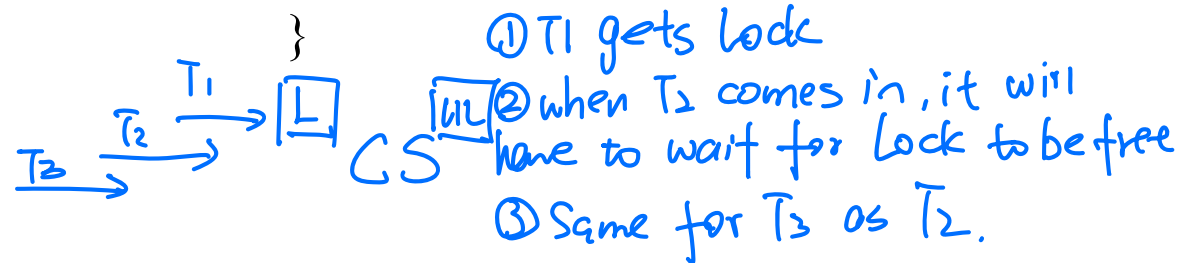
Mutex_Lock(L) {
  Disable Interrupts/Use Spinlock
  Check if anyone is using L
  If not {
    Set L to being used
  }
  else {
    Move this TCB to Blocked
    queue for L
    Select another activity to run
    from Ready queue
    Context switch to that activity
  }
  Enable Interrupts/Use Spinlock
}

```

```

Mutex_Unlock(L) {
  Disable Interrupts/Use spinlock
  if (blocked queue of L == NULL)
    Set L to free
  }
  else {
    Move TCB from head of
    Blocked queue of L to
    Ready queue
  }
  Enable Interrupts/Use spinlock
}

```



NOTE: These are OS system calls (where Disable/Enable are available), or Library calls (where Spinlocks are only option for implementation)

Until now ...

- Exclusion synchronization/constraint
 - Typical construct mutual exclusion lock
 - **Mutex_lock(m)** *different c.s. have different Lock(m;)*
 - **Mutex_unlock(m)**
 - Do a “man” on **pthread_mutex_lock()** for further syntactic/semantic information.

- **But you also need synchronization constructs for other tasks than exclusion (i.e., **ordering**)**
 - **E.g., If printer queue is full, I need to wait until there is at least 1 empty slot**
 - **Note that `mutex_lock()/mutex_unlock()` are not very suitable to implement such synchronization**
 - **We need constructs to enforce orderings (e.g., A should be done after B).**

enable a particular thread
put itself into a queue to wait for a particular condition
to be made, time

Condition Variables

→ inform one member of a block queue for that particular condition
variable that is now unblocked and move to ready queue

- C_wait() and c_signal() operations.
- A thread blocked on c_wait() returns when another performs a c_signal().

```
C_wait(C) {
    Disable Interrupts /Use spinlock
    Move this TCB to Blocked queue for C
    Select another thread to run
        from Ready queue
    Enable Interrupts /Use spinlock
    Context switch to that thread
}
```

```
C_signal(C) {
    Disable Interrupts/Use spinlock
    Check if blocked queue
        for L is empty
    else {
        Move TCB from head of
            Blocked queue of C to
            Ready queue
    }
    Enable Interrupts/Use spinlock
}
```

```
Cond_t not_full, not_empty;;
```

```
Int count == 0;
```

```
Append() {
```

```
    if count == N    c_wait(not_full);
```

*wait until
↑ buffer is not full*

*→ have to run atomic
to avoid rescheduling*

... ADD TO BUFFER, UPDATE COUNT ...

```
    c_signal(not_empty);
```

```
}
```

```
Remove() {
```

```
    if count == 0    c_wait(not_empty);
```

*wait until
↑ buffer is not empty*

... REMOVE FROM BUFFER, UPDATE COUNT

```
    c_signal(not_full);
```

```
}
```

However, there is something wrong with this code!
There is a gap between checking (count == N) and c_wait()!
Similarly, for Remove.

Solution: Put `c_wait()` within a `mutex_lock()`

Cond_t not_full, not_empty;

Mutex_lock m;

Int count == 0;

Append() {

mutex_lock(m);

if count == N c_wait(not_full,m);

... ADD TO BUFFER, UPDATE COUNT ...

c_signal(not_empty);

mutex_unlock(m);

}

Remove() {

mutex_lock(m);

if count == 0 c_wait(not_empty,m);

... REMOVE FROM BUFFER, UPDATE COUNT

c_signal(not_full);

mutex_unlock(m);

}

**C_wait(c,m) : You give up “m” before waiting, and you regain “m”
when you are signaled.**

Issue

- **But, this “solution” does not really work if there are two threads that run “remove” (or “append”)**
 - **Can you identify why not?**

Solution: Put `c_wait()` within a `mutex_lock()`

```
Cond_t not_full, not_empty;
```

```
Mutex_lock m;
```

```
Int count == 0;
```

```
Append() {
```

```
    mutex_lock(m);
```

```
    if count == N    c_wait(not_full,m);
```

```
    ... ADD TO BUFFER, UPDATE COUNT ...
```

```
    c_signal(not_empty);
```

```
    mutex_unlock(m);
```

```
}
```

```
Remove() {
```

```
    mutex_lock(m);
```

```
    if count == 0    c_wait(not_empty,m);
```

```
    ... REMOVE FROM BUFFER, UPDATE COUNT
```

```
    c_signal(not_full);
```

```
    mutex_unlock(m);
```

```
}
```

Suppose that one thread performs a “`c_wait`”. But, another thread runs “`remove`” after a “`c_signal`” and **steals** the buffer element.

Textbook

- **The textbook goes through these scenarios in detail in Chapter 30**
 - **Need to understand these for P2**

Solution: Put `c_wait()` within a `mutex_lock()`

`Cond_t not_full, not_empty;`

`Mutex_lock m;`

`Int count == 0;`

`Append() {`

`mutex_lock(m);`

`while count == N c_wait(not_full,m);`

`... ADD TO BUFFER, UPDATE COUNT ...`

`c_signal(not_empty);`

`mutex_unlock(m);`

`}`

`Remove() {`

`mutex_lock(m);`

`while count == 0 c_wait(not_empty,m);`

`... REMOVE FROM BUFFER, UPDATE COUNT`

`c_signal(not_full);`

`mutex_unlock(m);`

`}`

Need to check that the condition still holds when a thread is finally scheduled.

Broadcast

- What if you do not know which thread will satisfy the condition when exiting a critical section?
 - I.e., you may “signal” a thread that does not meet the condition?
- What happens?
 - It goes back to waiting.
 - And a thread that meets the condition may not be awoken
- **pthread_cond_broadcast** - shall unblock all threads currently blocked on the specified condition variable *cond*

