

# Concurrency Control via Synchronization



# Need for Synchronization

- Activities share resources.
- It is important to **coordinate** their progress to ensure proper usage.

# Examples of coordination

- John and Mary are each printing a different 10 page document on the same printer. You do not want their pages/output interleaving!

– Exclusion

- John and Mary are sharing a bank account. John deposits \$10. Mary should be allowed to withdraw only after this deposit

- Ordering


# Resources

- **There are different kinds of resources that are shared:**
  - Physical (terminal, disk, network, ...)
  - Logical (files, sockets, memory, ...)
- **For the purposes of this discussion, let us focus on “memory” to be the shared resource**
  - **i.e. threads can all read and write into memory (variables) that are shared.**


# Problems due to sharing

- Consider a shared printer queue,  
`spool_queue[N]`
- 2 threads want to enqueue an element each to this queue.
- `tail` points to the current end of the queue
- Each thread needs to do *Thread will look at tail to know where to put element*  
`tail = tail + 1;`  
`spool_queue[tail] = "element";`

# What we are trying to do ...



# What is the problem?

- $\text{tail} = \text{tail} + 1$  is NOT 1 machine instruction
- It can translate as follows:
  - Load tail, R1
  - Add R1, 1, R2
  - Store R2, tail
- These 3 machine instructions may NOT be executed atomically. 

*atomically*




- **If each thread is executing this set of 3 instructions, context switching can happen at any time.**
- **Let us say we get the following resultant sequence of instructions being executed:**
  - 1: Load tail, R1**
  - 2: Load tail, R1**
  - 1: Add R1, 1, R2**
  - 2: Add R1, 1, R2**
  - 1: Store R2, tail**
  - 2: Store R2, tail**

# Leading to ...

## Thread 1

$\text{tail} = \text{tail} + 1;$   
 $\text{Spool\_queue}[\text{tail}] = X$



## Thread 2


$\text{tail} = \text{tail} + 1;$   
 $\text{Spool\_queue}[\text{tail}] = Y$



race condition: result of a sequence of instructions depends on its interleaving with other instructions.

- Situations like this that can lead to erroneous execution depending on who executes when, are called **race conditions**.  
happen in shared memory space like global variables
- Debugging race conditions is NOT easy! since errors can be non-repeatable.

# Avoiding Race Conditions

- If we had a way of making those (3) instructions atomic – i.e., while one thread is executing those instructions, another cannot execute the same instructions – then we could have avoided the race condition.
- These 3 instructions are said to constitute a **critical section**. 
- While one thread is in a critical section, another should NOT be allowed to execute the same critical section – **mutual exclusion**. 