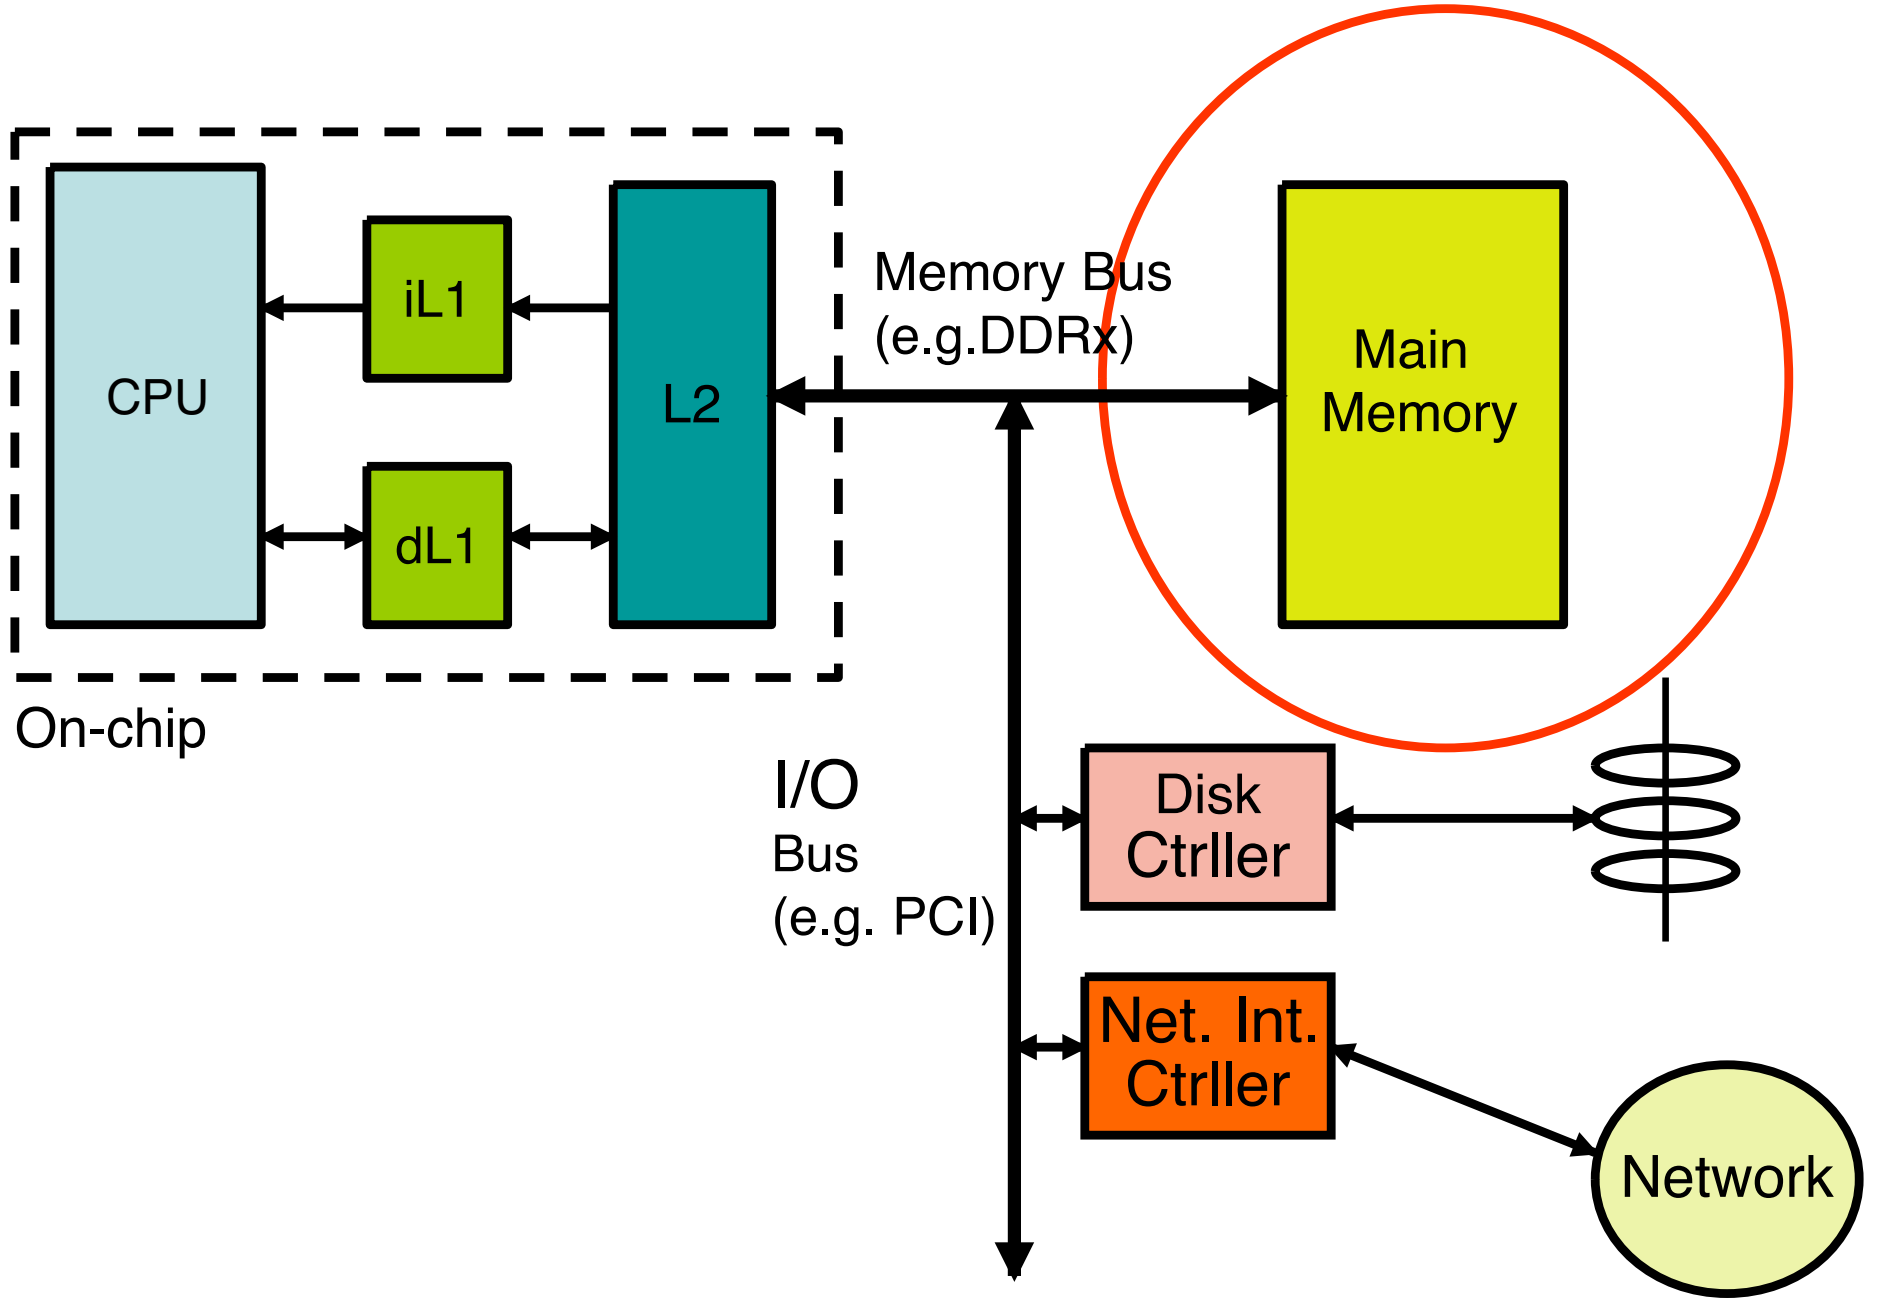
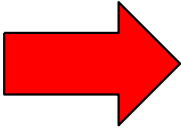


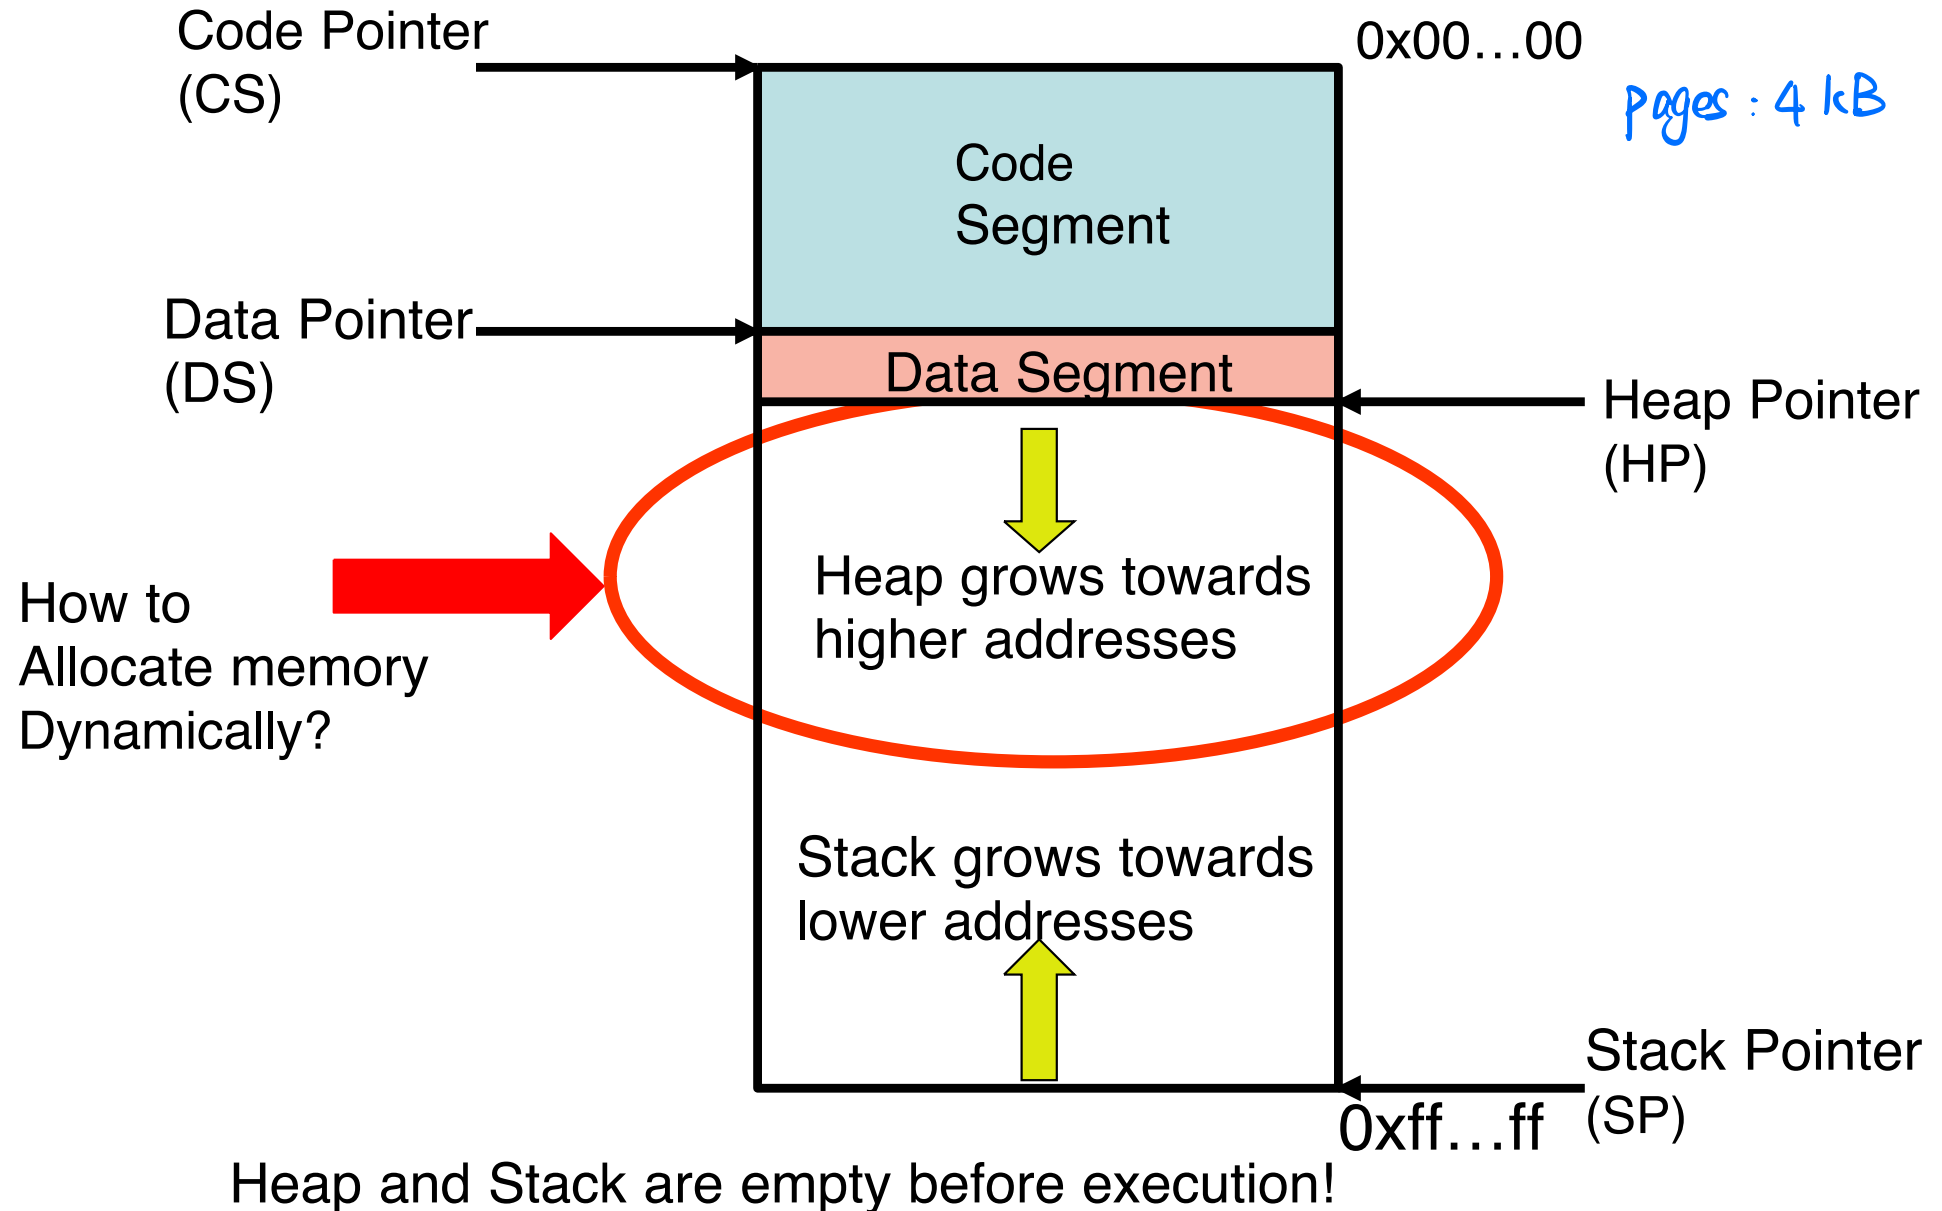
Memory Management



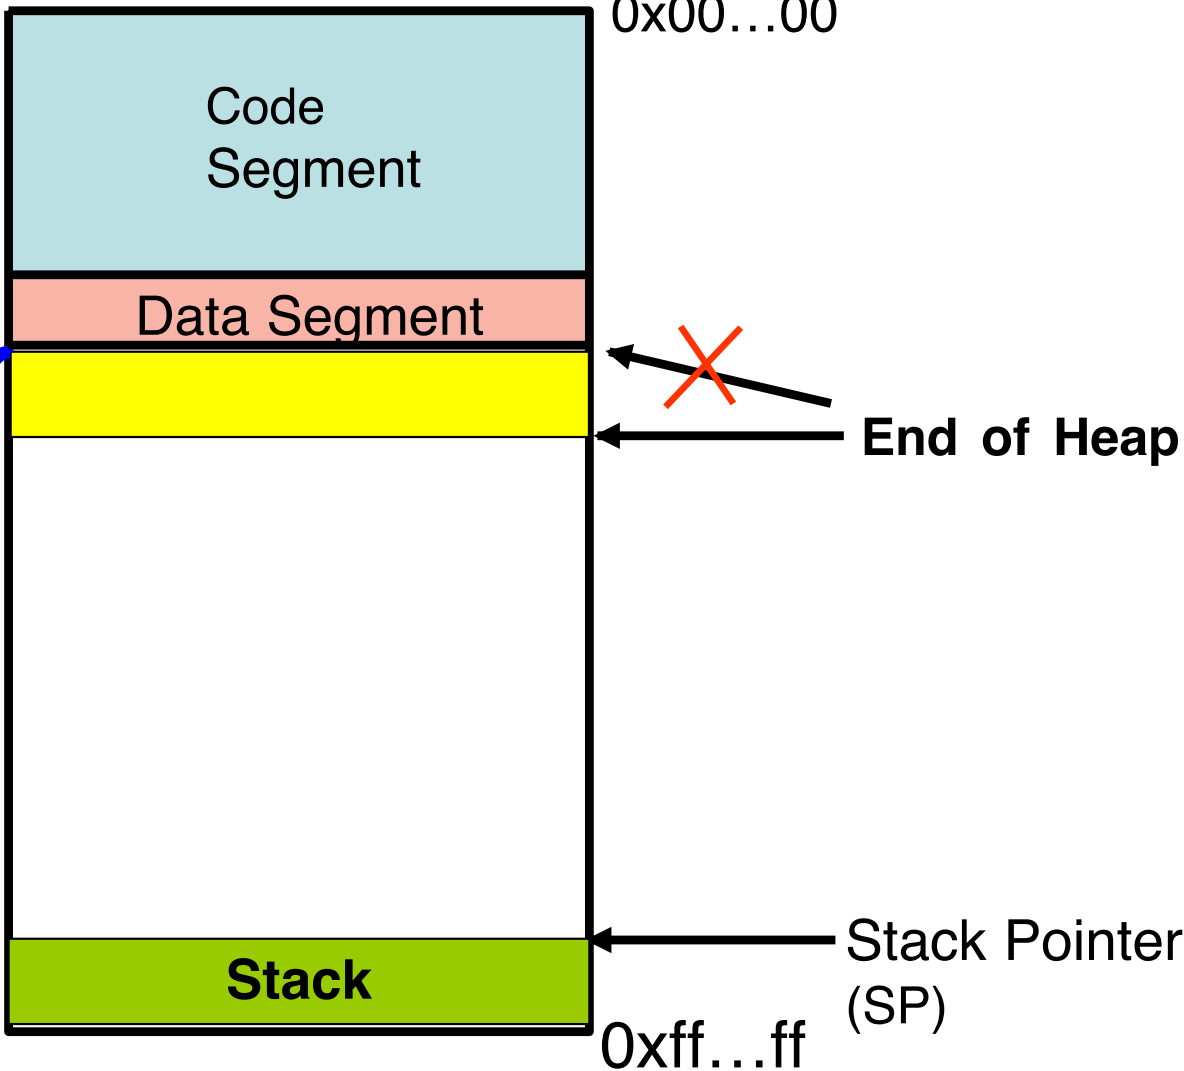

Need for Memory Management

- **Spatial issues in Allocation**
 -  – How to implement alloc/free?
- **Physical memory (DRAM) is limited**
 - A single process may not all fit in memory
 - Several processes (their address spaces) also need to fit at the same time
- **Disallow 1 process from accessing another processes memory.**

Address Space of a Process



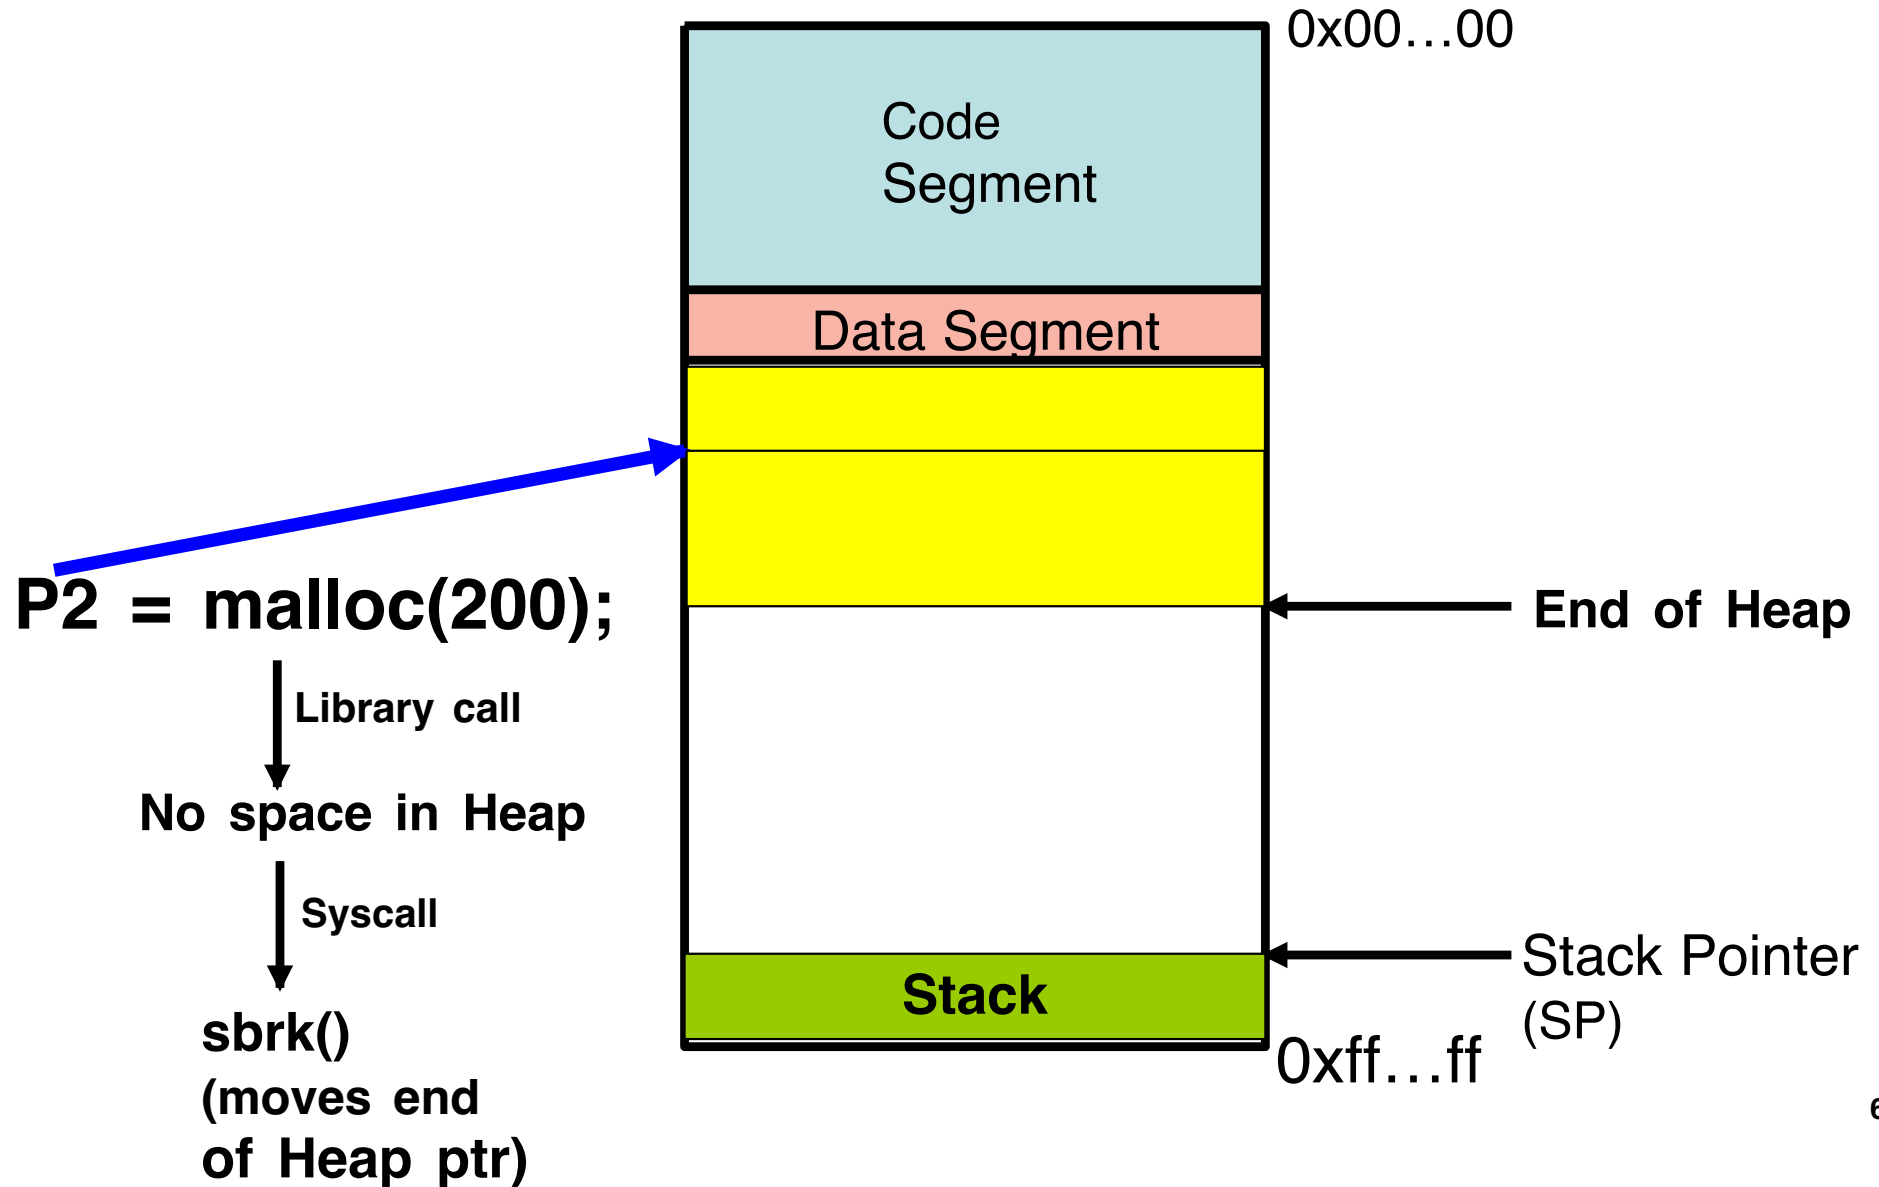
Malloc()



P1 = malloc(100);

Library call
↓
No space in Heap
↓
Syscall
↓
sbrk()
(moves end of Heap ptr)

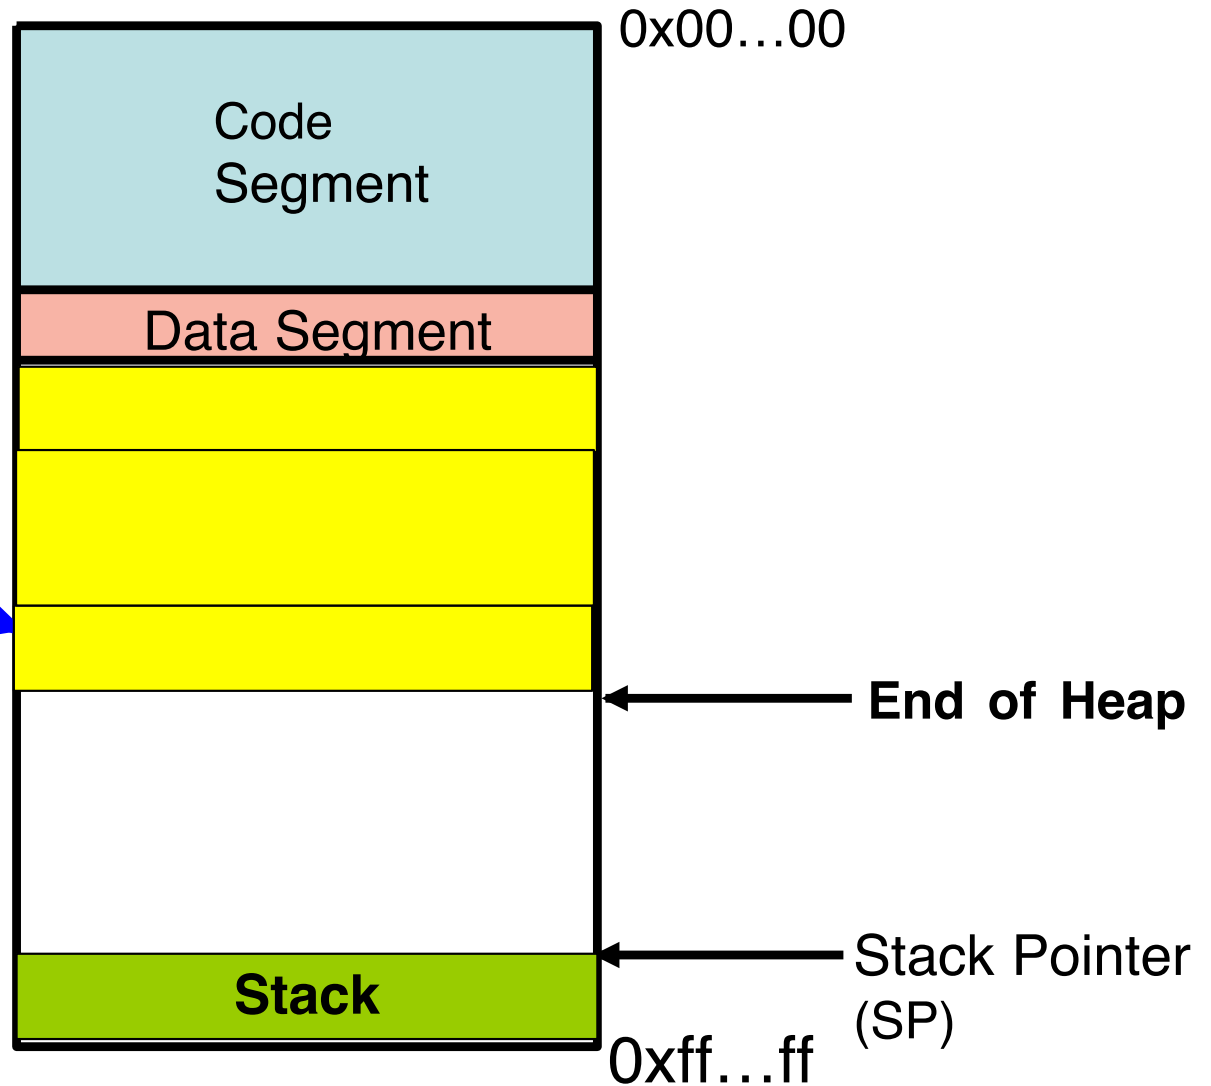
Malloc()



Malloc()

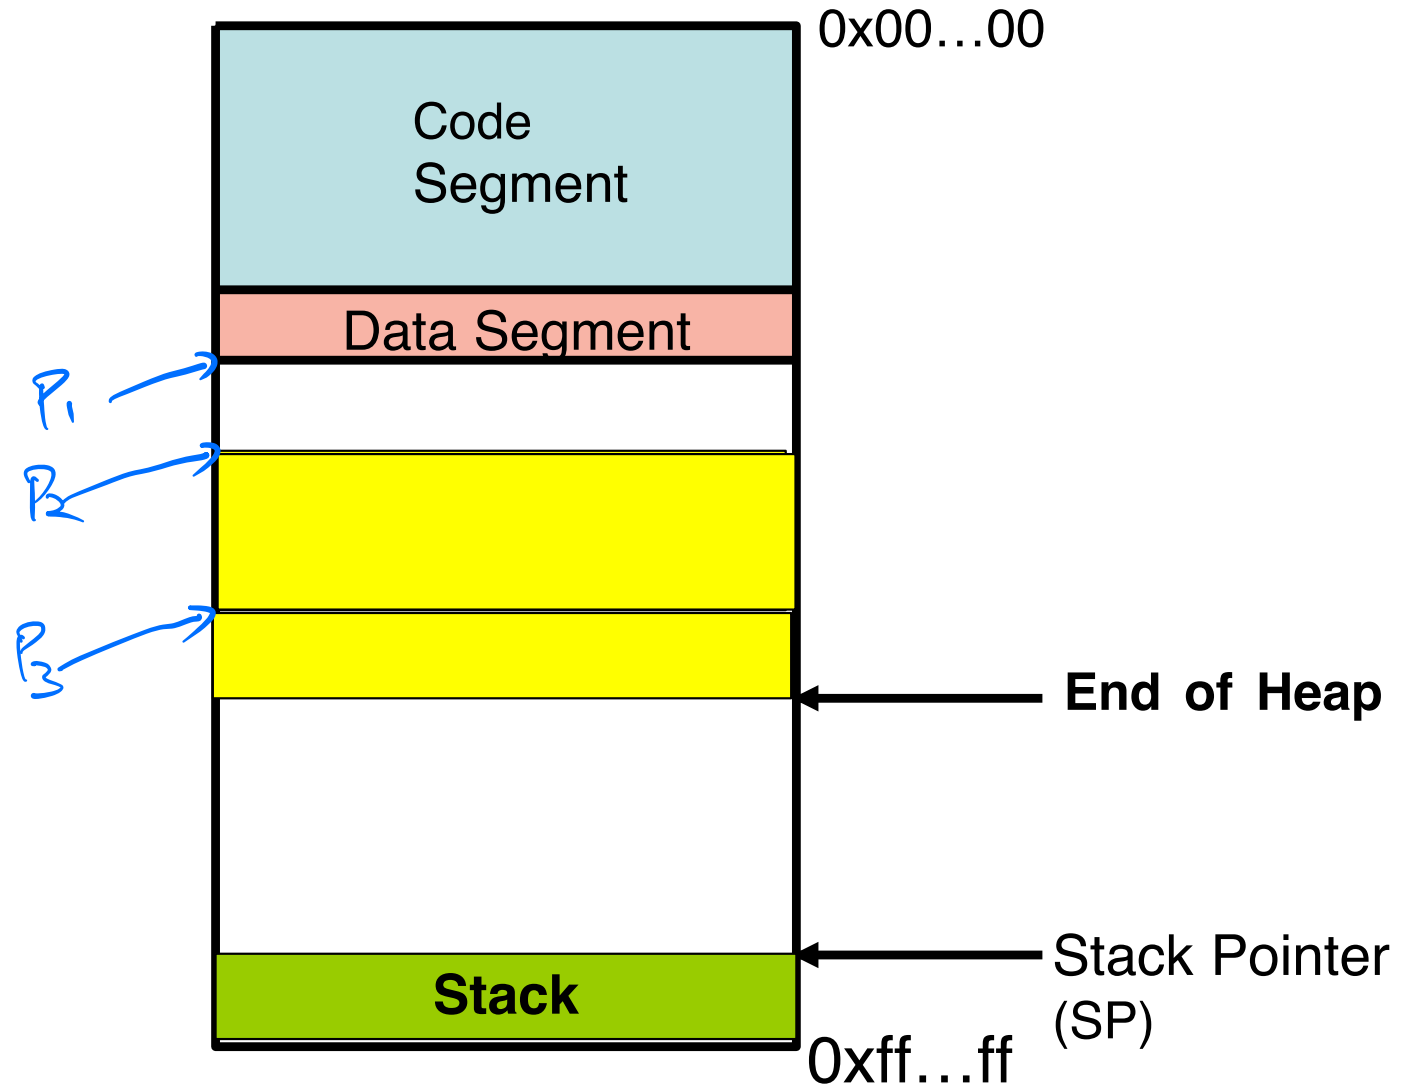
P3 = malloc(100);


Library call
 ↓
No space in Heap
 ↓
 Syscall
 ↓
sbrk()
 (moves end of Heap ptr)




Malloc()

free(p1);



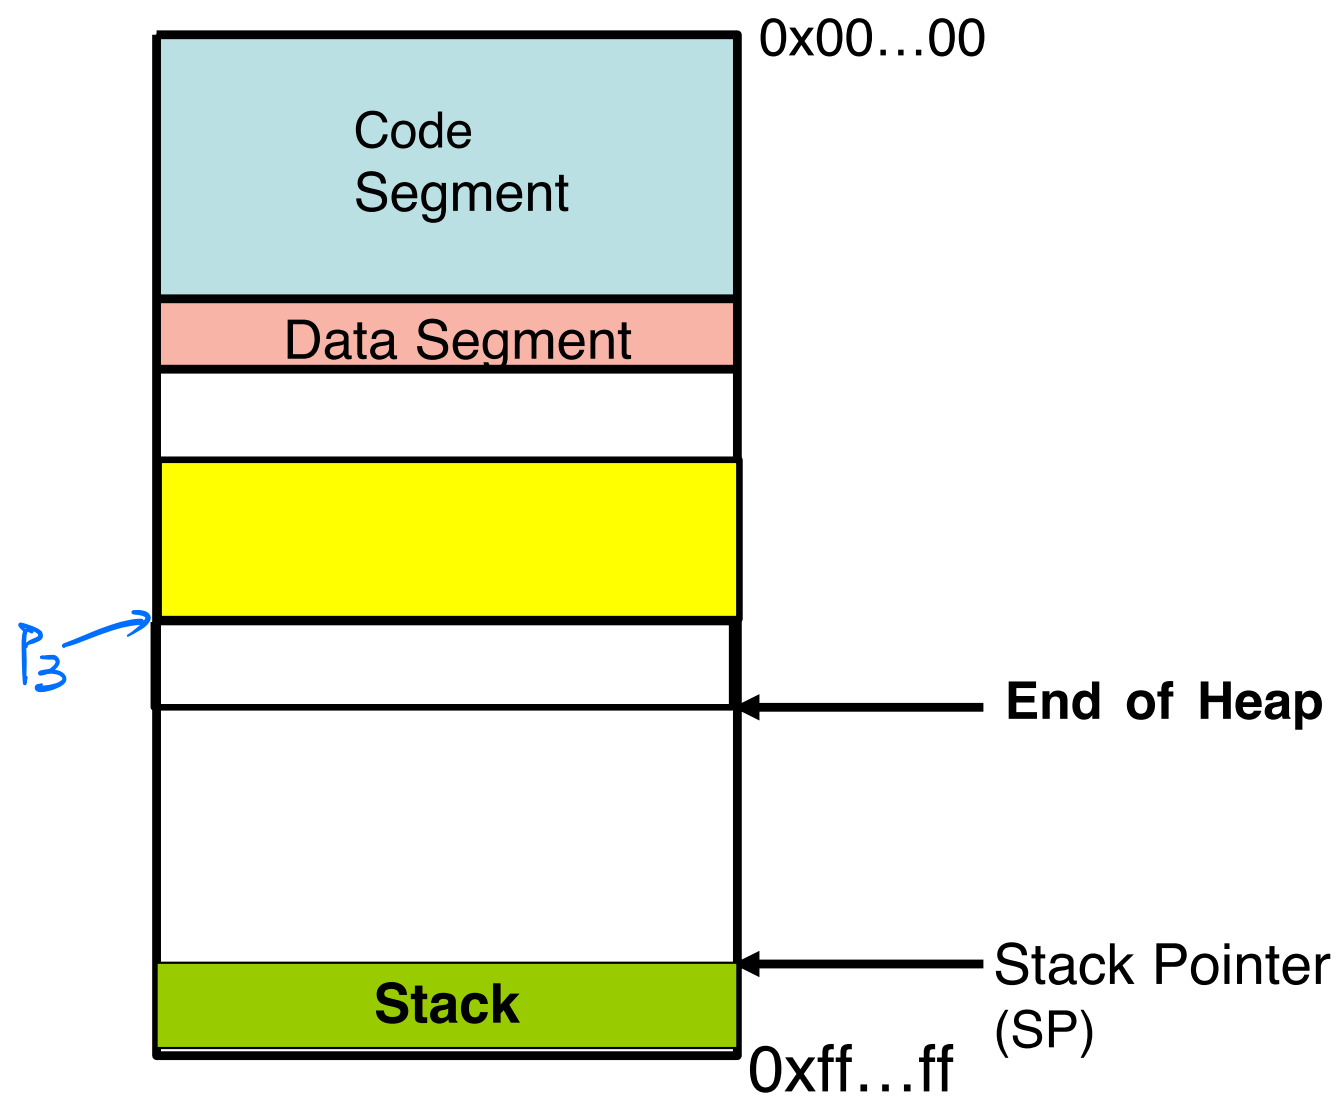


Physical



Malloc()

free(p3);

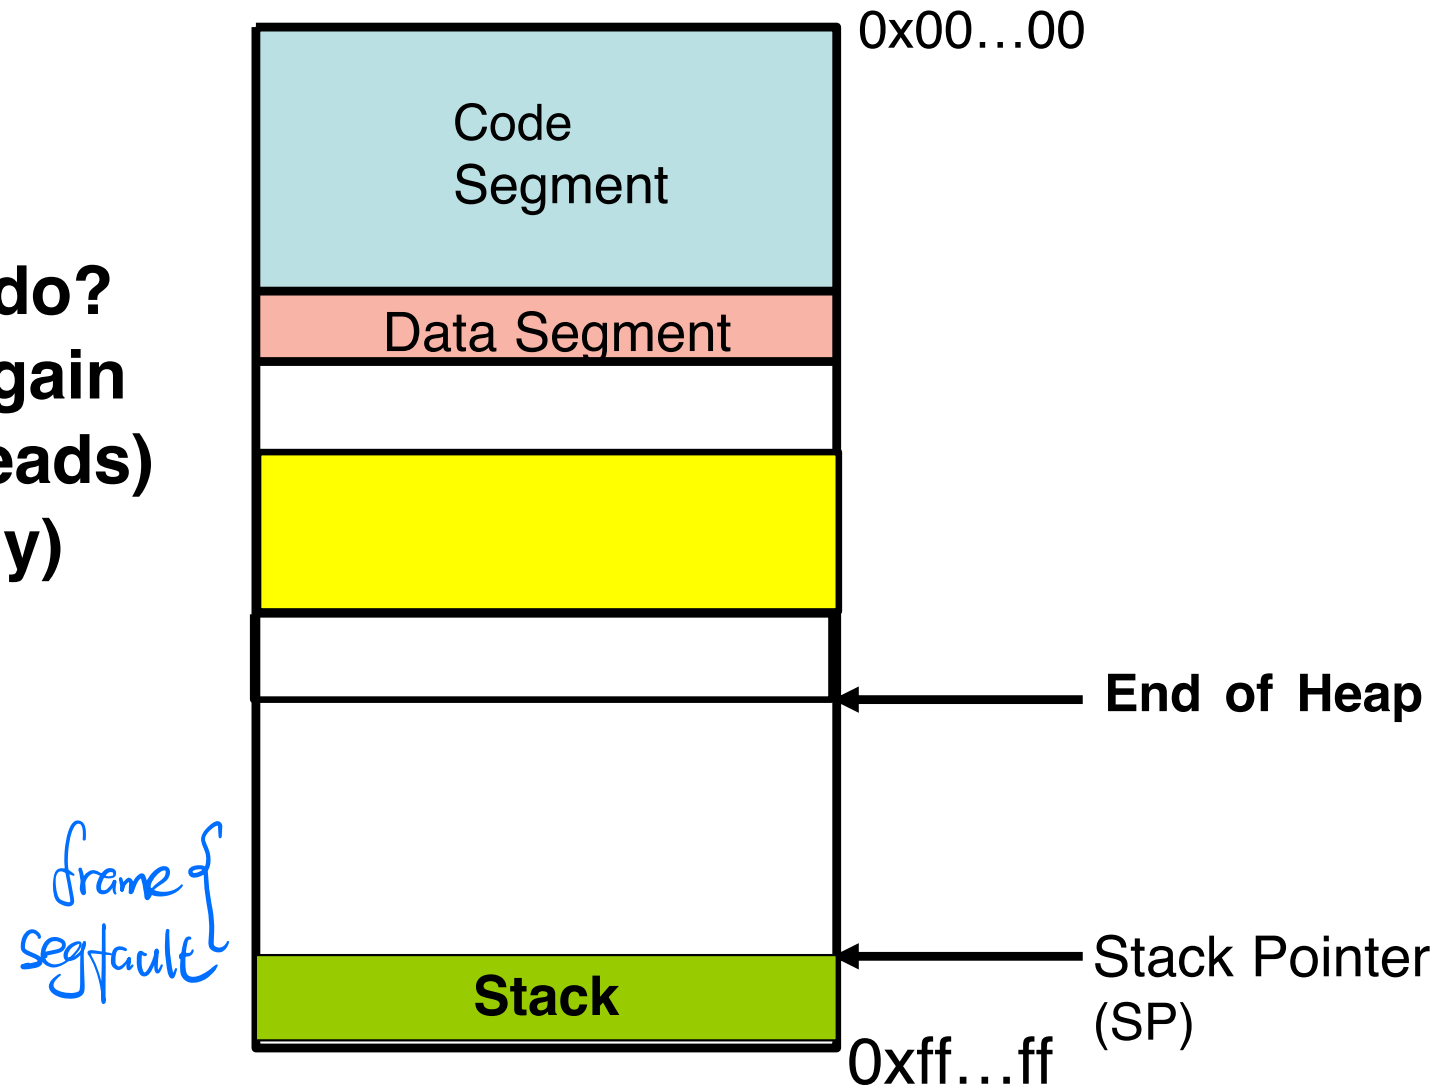


Malloc()

Malloc(150)


What should we do?

- **Extend heap again (syscall overheads)**
- **Compact (costly)**



Makes allocation of free space to malloc an important problem


The Memory Allocation Problem



Question: How do we perform these allocations?

Goals

- Program wants allocation to be contiguous. Note “`p = malloc(1000)`” requires 1000 bytes to be contiguous for `p[842]` to work



- `Allocation()` and `Free()` should be fairly efficient
- Should be able to satisfy more requests at any time (i.e., the sum total of holes should be close to 0 with waiting requests).

Solution strategies


- **Contiguous allocation**
 - The requested size is granted as 1 big contiguous chunk.
 - E.g., first-fit, best-fit, worst-fit, buddy-system *– Algorithm*
- **Non-contiguous** physically but virtually contiguous
 - The requested size is granted as several pieces (and typically each of these pieces is of the same – fixed - size).
 - E.g., **paging**
- Use the former to do virtually contiguous in latter

Contiguous allocation

- **Data structures:**
 - List of allocated regions
 - List of holes.
- Find a hole and make the allocation (and it may result in a smaller hole).
- Eventually, you may get a lot of holes that become small enough that they cannot be allocated individually.
- This is called **external fragmentation**.

external

Easing external fragmentation



Note that this can be done only with relocatable code and data (use indirect/indexed/relative addressing)

But compaction is expensive and we want to do this as infrequently as possible.

Contiguous allocation

- Which hole to allocate for a given request?
- **First-fit**
 - Search through the list of holes. Pick the first one that is large enough to accommodate this request.
 - Though allocation may be easy, it may not be very efficient in terms of fragmentation.

- **Best Fit**

- Search through the entire list to find the smallest hole that can accommodate the given request.
- Requires searching through the entire list (or keeping it in sorted order).
- This can actually result in very small sized holes making it undesirable.

- **Worst fit**


- Pick the largest hole and break that.
- The goal is to keep the size of holes as large as possible.
- Allocation is again quite expensive (searching through entire list or keeping it sorted).

What do you do for a `free(p)`?

- You need to determine whether nearby regions (on either side) is free, and if so you need to make a larger hole.
- To do this:
 - We need to know what size you are free-ing
 - We need to know whether either side of this is free

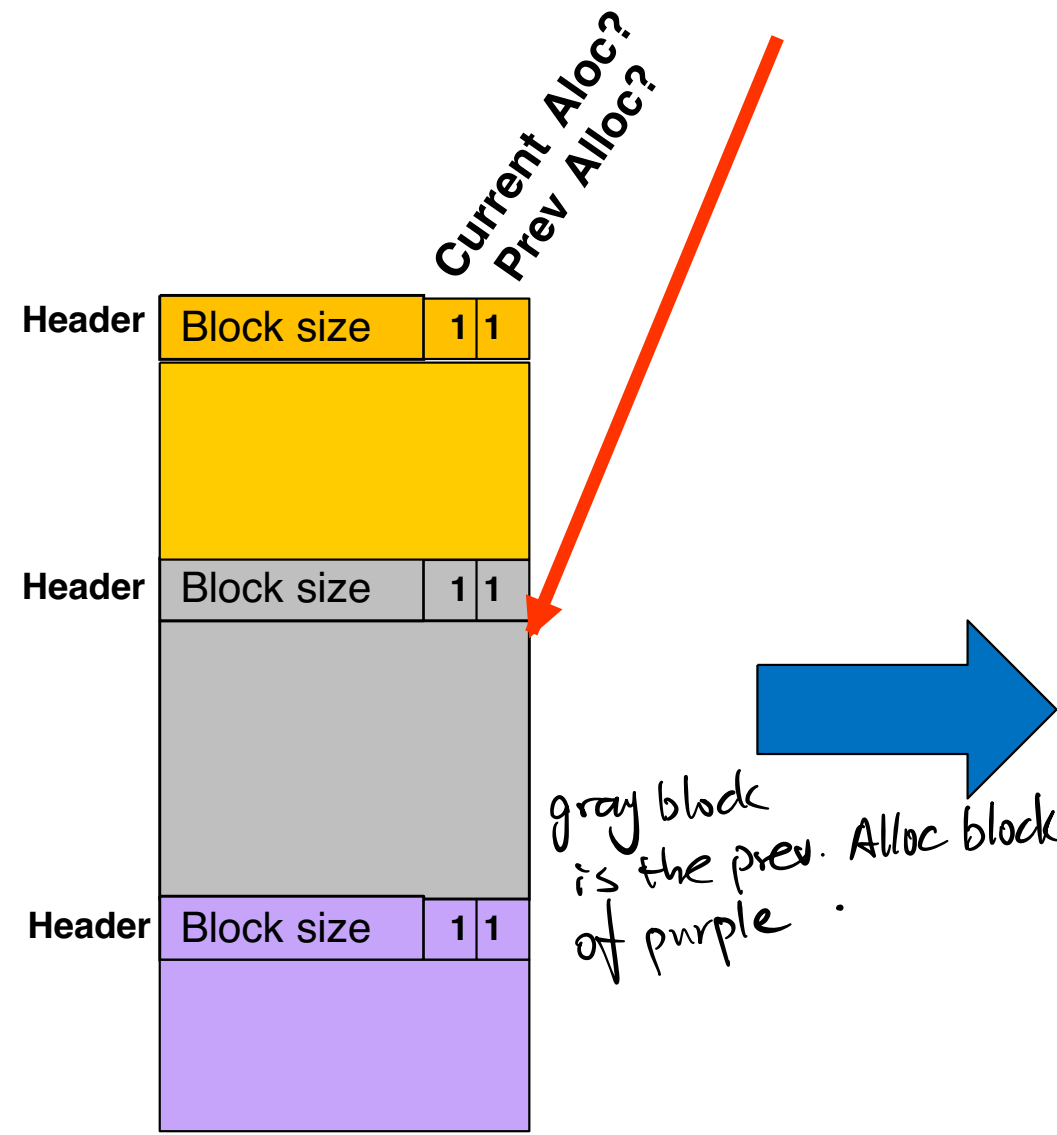
How do you know what size you are freeing?

- **Allocate extra space for a small header and return pointer after this header for an allocation.**
- **Do a negative offset of the pointer passed by free to find size!**




How do you know whether either side of what you are free-ing are free, and if so merge with them to create a larger hole?

free(p) – Case 1

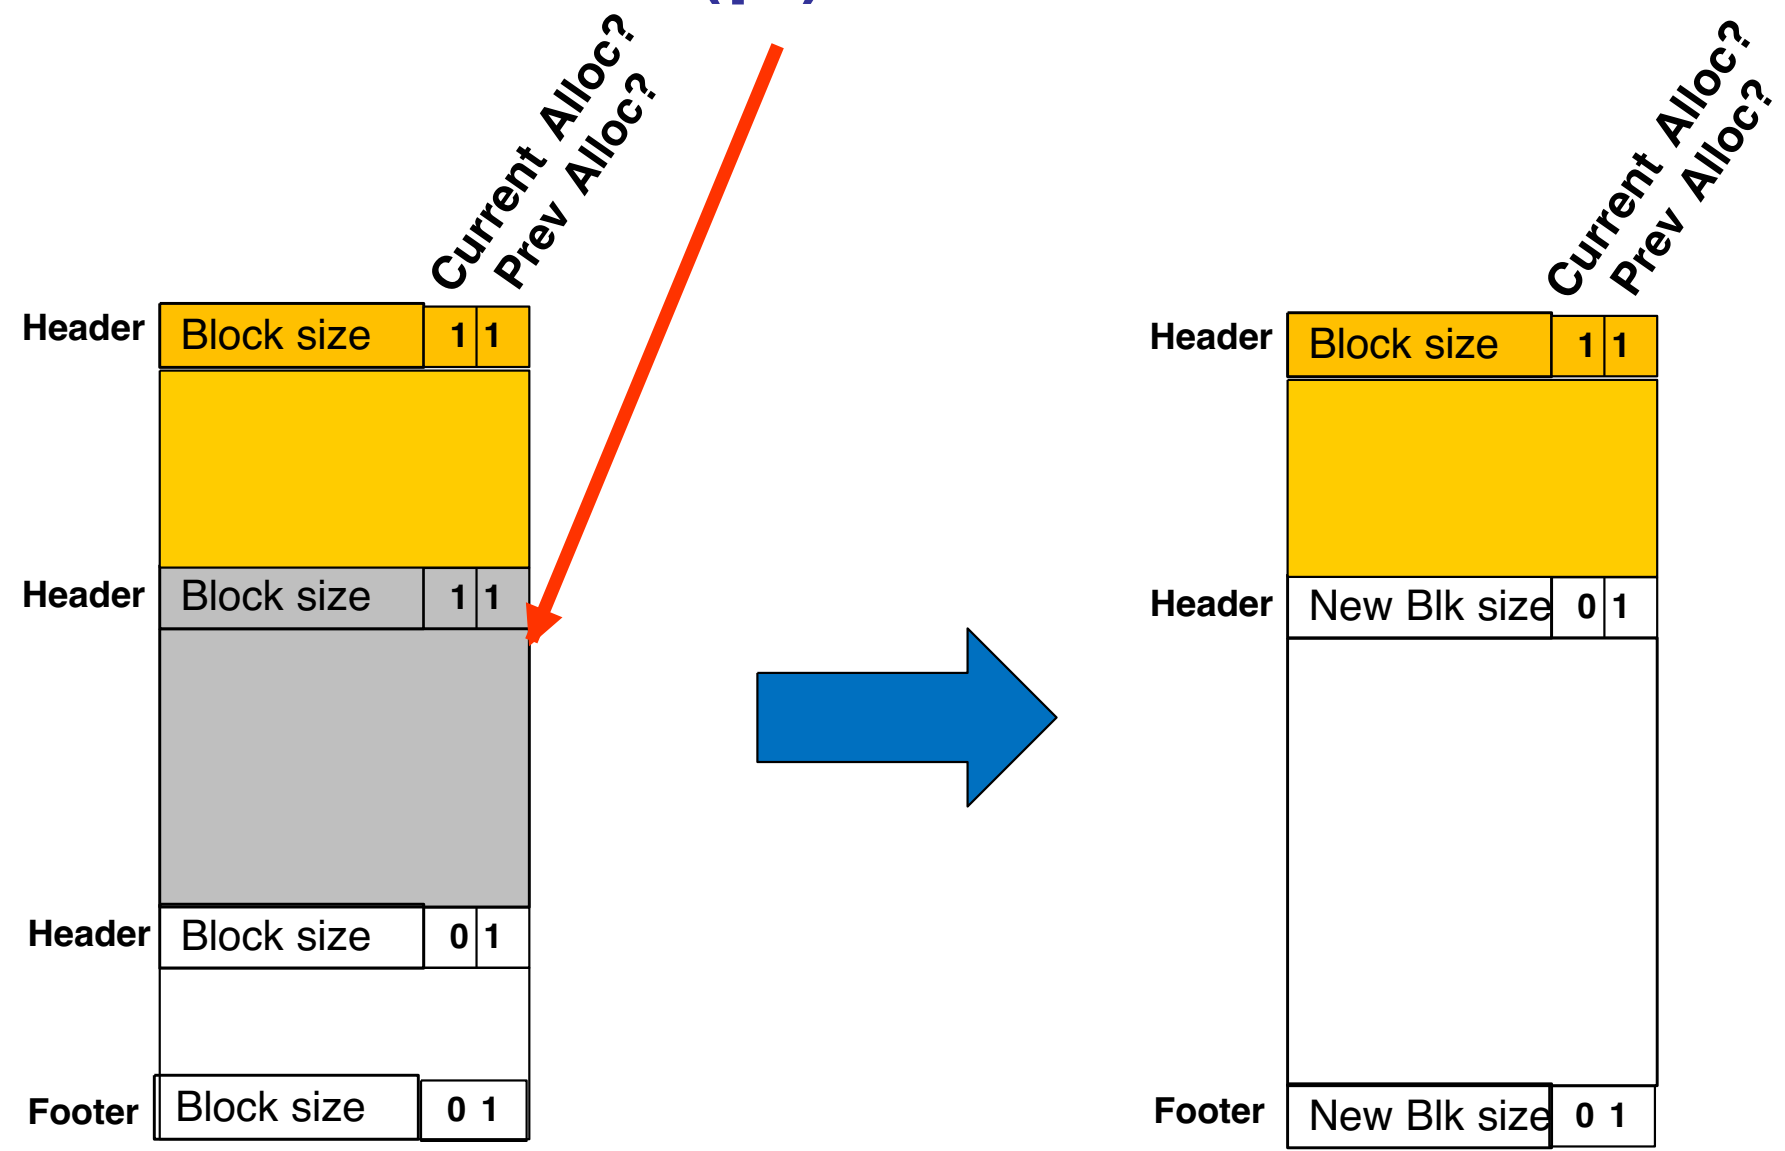


gray block
is the prev. Alloc block
of purple.

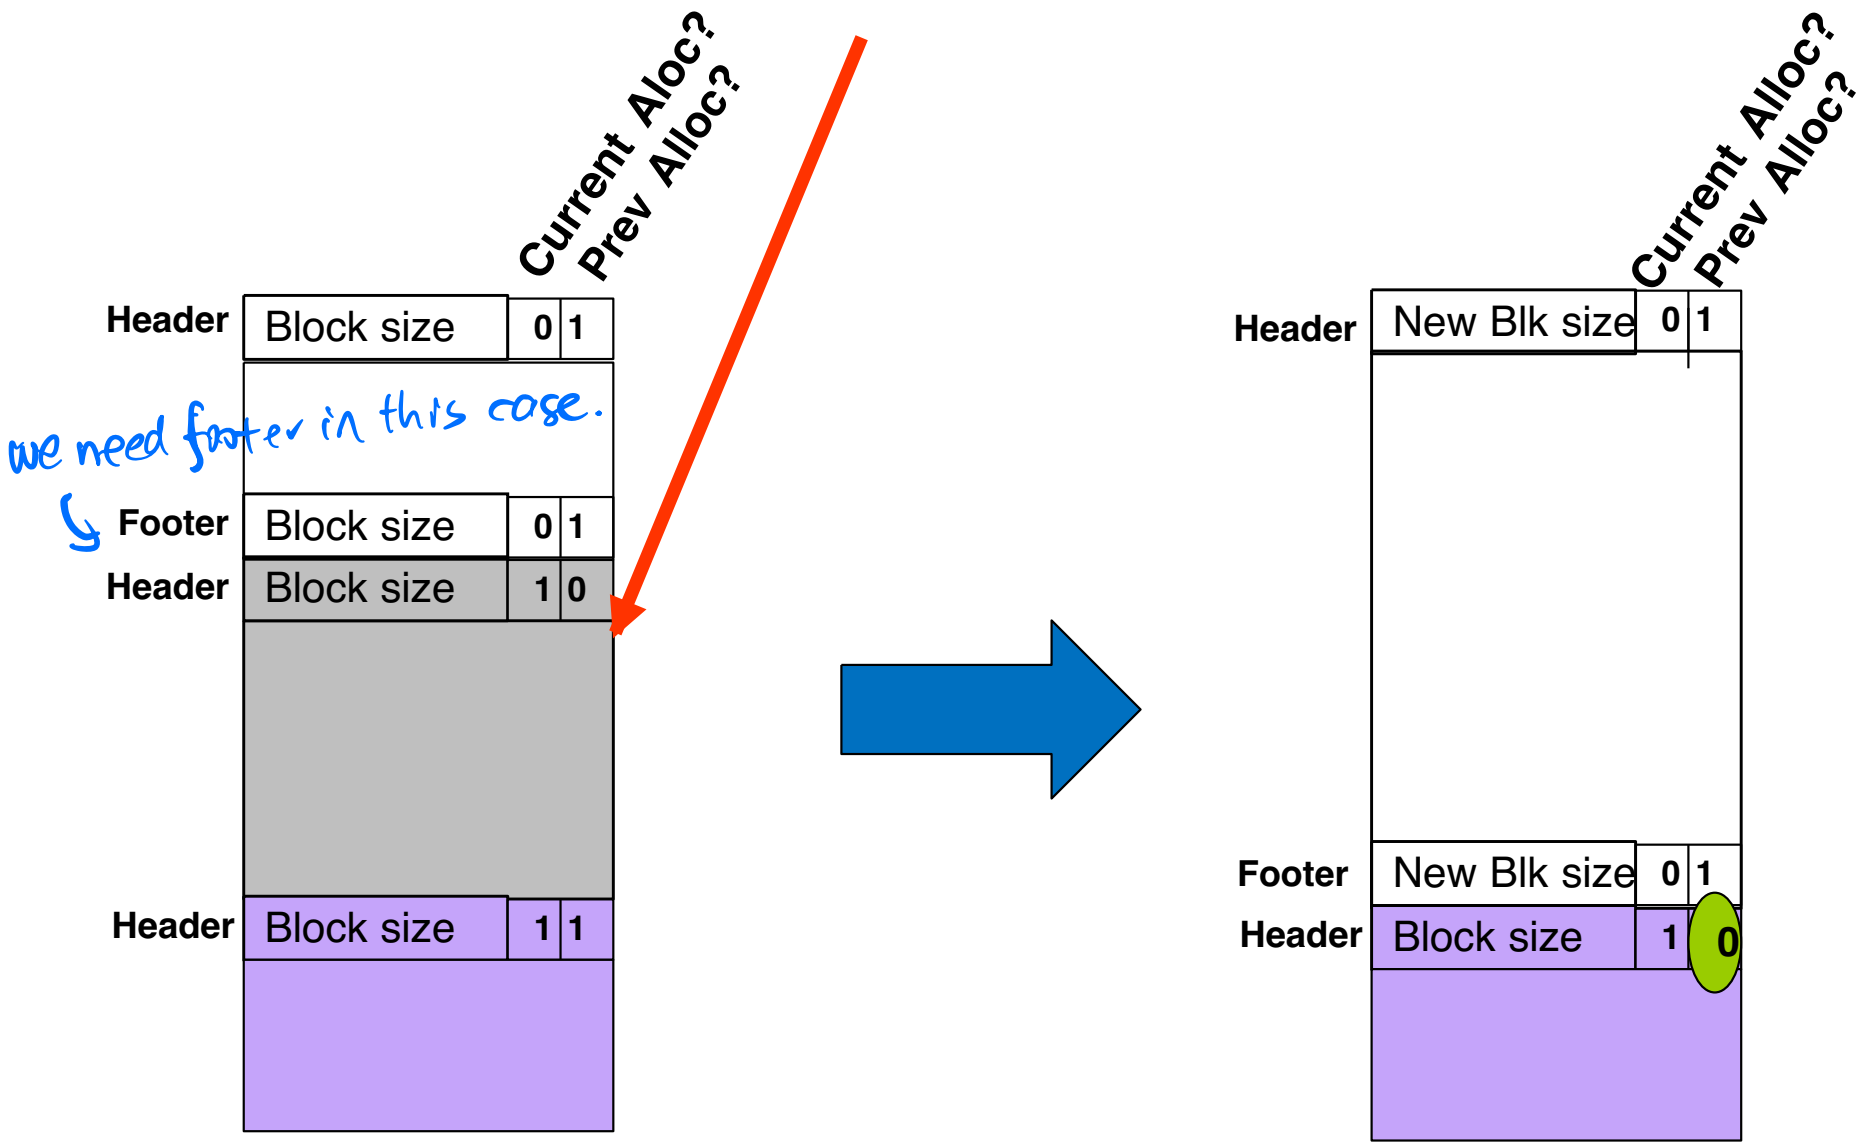


free(ptr)
header->next = ptr + header
⇒ blk-size + size of ²²
(footer type)


free(p) – Case 2



free(p) – Case 3



free(p) – Case 4



Costs

- **Free: Look only at neighbors ($O(1)$)**
- **Allocation: Search through list ($O(N)$).**

Can we do better?

Buddy System



- **$\text{Log}(N)$ cost for allocation/free.**
- **Round up request to next power of 2, and find a block to allocate**

An Example




Request

256K




Release

B




Release

A



Request


75K






Release

C




Release

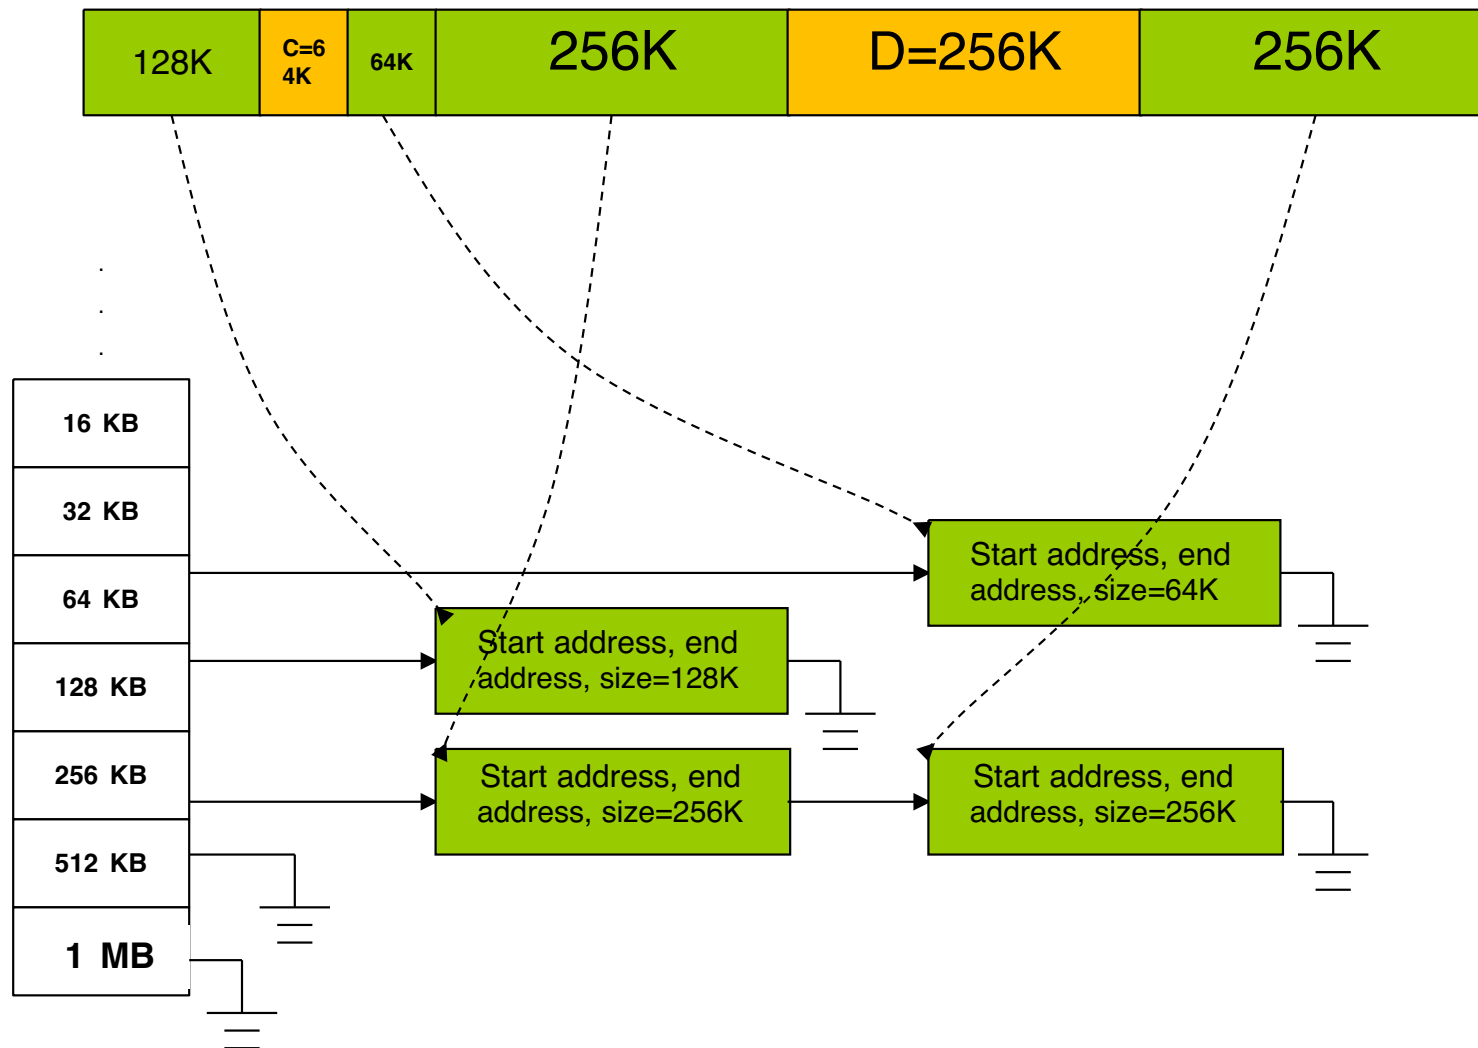
E



Release

D






List of Available Holes

Problems with Buddy

—might waste memory



```
Object *objs[4];
```

```
for (i=0; i < 4; i++){
```

```
    obj[i] = malloc(600)
```

```
}
```

Object size is 600 Bytes.

The problem can worsen for larger sizes because of rounding off to next power of 2 !


This is called Internal Fragmentation!

- **Can we leverage such repetitive behavior?**
- **Can we pack more objects into contiguous chunks to avoid internal fragmentation?**

Slab Allocator


- Do continuous allocation of large chunks (e.g. 4K, 8K, ...) – called **Slabs** – using *Buddy* allocator.
- Each Slab will contain objects of same size (e.g. 600 bytes in this example) – **Slab type**
- An object of a particular size will try to get allocated in the slab of that type/size.
- **Slab size** is the number of objects of that type/size that this slab can hold.
- When you try to allocate more objects than what a single slab can hold, you allocate one more slab for that type (using Buddy).

Slab Allocator



Slab Descriptor Table				
Type	Size	#Obj	#Used	Slab-ptr
600	4K	6	4	

```
SmallObject* sm_objs[4];
for (int i=0; i < 4; i++) {
    sm_objs[i] = malloc(600);
}
```

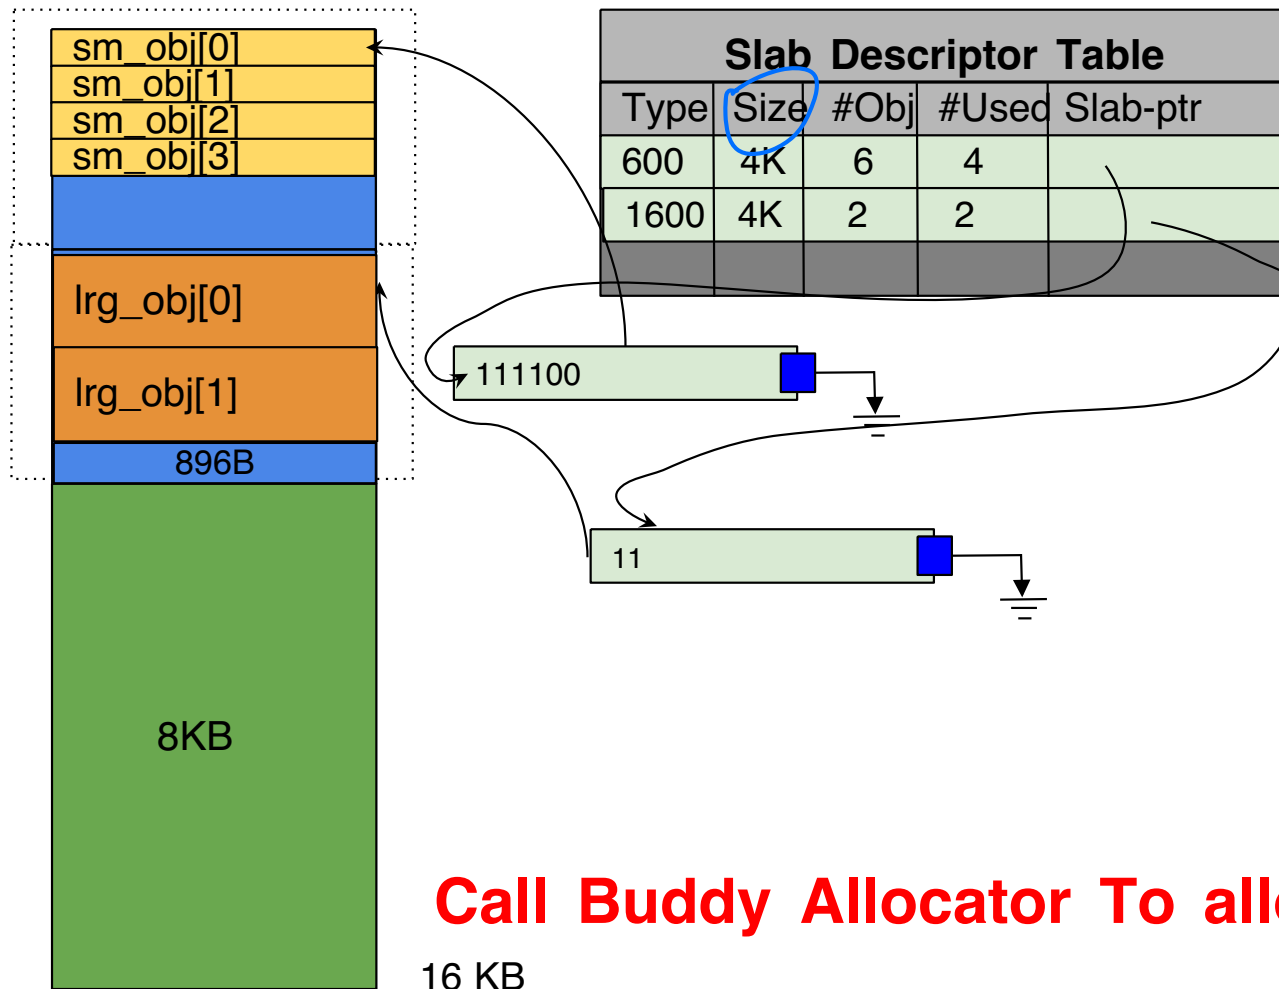


Bitmap of what is allocated and what is free in this slab.

Call Buddy Allocator To allocate new Slab

16 KB

Slab Allocator



```
SmallObject* sm_objs[4];
for (int i=0; i < 4; i++) {
    sm_objs[i] = malloc(600);
}
```

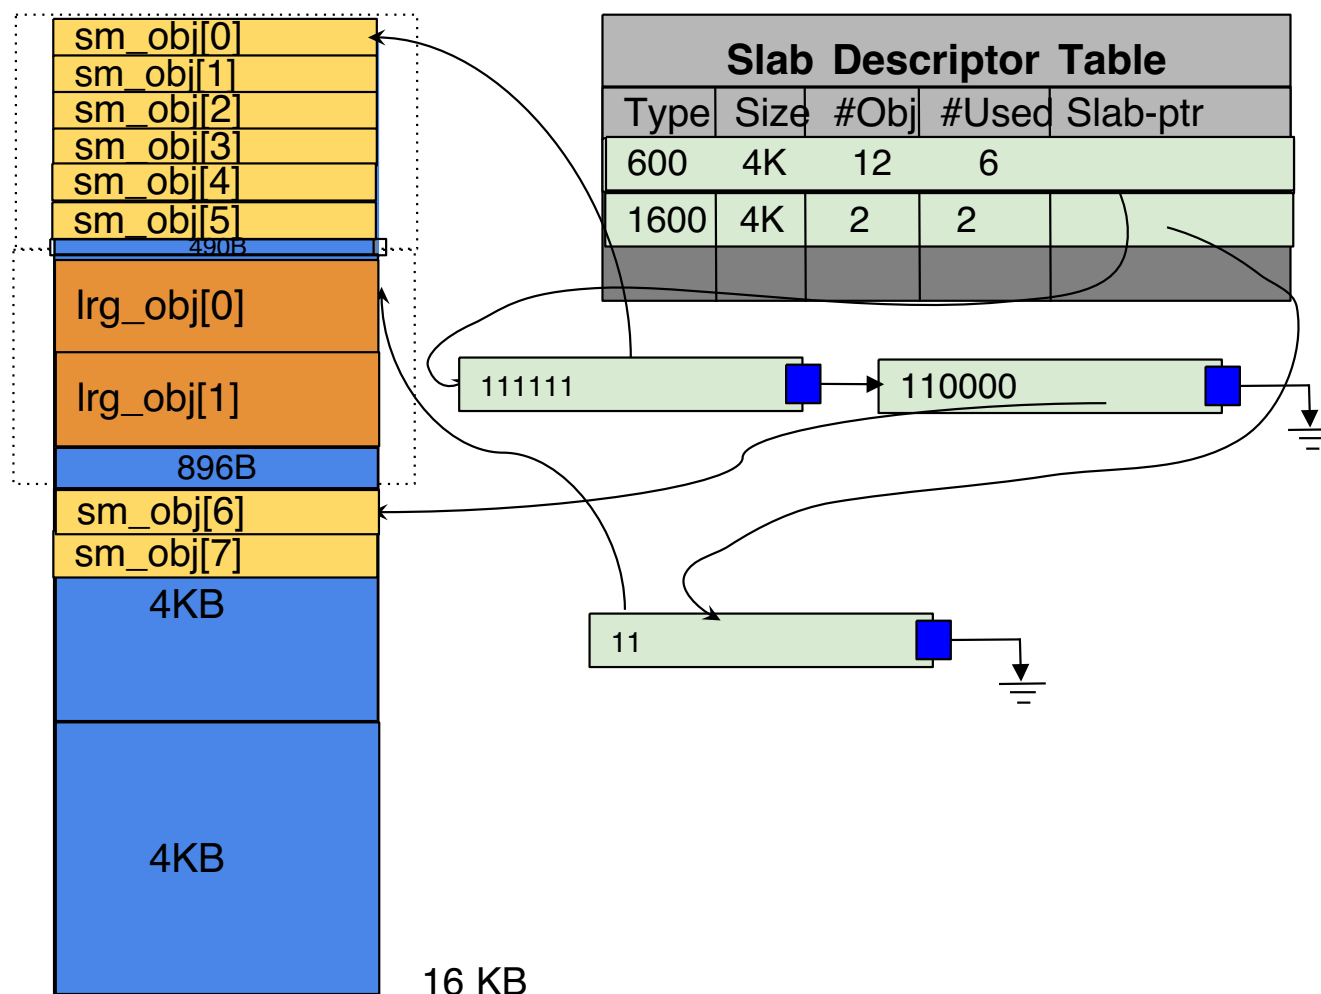
Larger object

```
LargeObject* lrg_objs[2];
for (int i=0; i < 2; i++) {
    lrg_objs[i] = malloc(1600);
}
```

Call Buddy Allocator To allocate new Slab

16 KB

Slab Allocator



```
SmallObject* sm_objs[4];
for (int i=0; i < 4; i++) {
    sm_objs[i] = malloc(600);
}
```

```
LargeObject* lrg_objs[2];
for (int i=0; i < 2; i++) {
    lrg_objs[i] = malloc(1600);
}
```

More small objects

```
SmallObject* sm_objs[4];
for (int i=0; i < 4; i++) {
    sm_objs[i] = malloc(600);
}
```

Call Buddy Allocator To allocate new Slab

Slab Allocator: Free

- Again use negative offset to find size of what is being freed.
- Use that to find the corresponding entry in Slab Descriptor Table.
- Go through one slab after another till you find where the chunk to be freed lies.
- Reset the “bit corresponding to that chunk” in the bitmap field.
- A future allocation of the same type can use this freed chunk