

# Processes

$PID=0$  , child

# Doing Well in This Course

- **What do you need to concentrate on to do well in this course?**



# What Is Hard?

- There are a key expectations we have for the students in this course
- **Terminology**: There is a lot of vocabulary in this course that you need to learn and use correctly – winging it won't work
- **Greedy Algorithms**: OSes solve a lot of computationally complex problems with approximate solutions – know them & tradeoffs
- **C Programming**: OSes use C language to manage memory extensively in ways that can create subtle errors – know pointers and debuggers

# Learning Terminology

- **Challenge:** Hard to know what is *really important* given a lot of potentially intimidating stuff - new to many
- What helps? Readings. Really!
- **Problem:** Reading about various OS details can be boring as it may be hard to know where this all goes or how it comes together
- What do you suggest?

# Learning Terminology

- **Challenge:** Hard to know what is really important given a lot of potentially intimidating stuff - new to many
- What helps? Readings. Really!
- **Problem:** Reading about various OS details can be boring as it may be hard to know where this all goes or how it comes together
- **Suggestion:** read/skim to pick out key concepts from the many pages, then read/study again in with those concepts in mind
- **Suggestion:** Take notes of things I emphasize from lectures (may not be on the <sup>5</sup>

# Learning Greedy

- **Challenge:** To solve a hard problem efficiently, people try lots of things and you need to know these options and their trade-offs
- What helps? Focus on the main problem
- **Problem:** There can be a lot of noise as we discuss this and that option – i.e., lose the forest for the trees
- What do you suggest?

# Learning Greedy

- **Challenge:** To solve a hard problem efficiently, people try lots of things and you need to know these options and their trade-offs
- What helps? Focus on the main problem
- **Problem:** There can be a lot of noise as we discuss this and that option – lose the forest for the trees
- **Suggestion:** Focus on how these solve the same problem differently
- **Suggestion:** And the resultant effects of these algorithm choices on operation

# Learning C

- **Challenge:** We are really going to use C in ways that leverage its power and danger in managing memory – beyond 311
- What helps? Mental model
- **Problem:** This may be the first language you have experiences with data objects and memory objects
- What do you suggest?



# Learning C

- **Challenge:** We are really going to use C in ways that leverage its power and danger in managing memory – beyond 311
- What helps? Mental model
- **Problem:** This may be the first language you have experiences with data objects and memory objects (pointers)
- **Suggestion:** Use the debugger to see how memory is represented and used (threads!)
- **Suggestion:** Learn safe programming techniques to avoid creating errors

# Topic for Today: Processes

# Program vs. Process

- What we looked at until now was a program (and its executable)
- It is not yet a process!
- **A process is a program in execution.**
- Think of a program as the recipe (instructions) for making a cake.
- The process is the “activity” of making the cake.
- A process has an associated program that it is executing and a state at any point of time.

# Program to Process

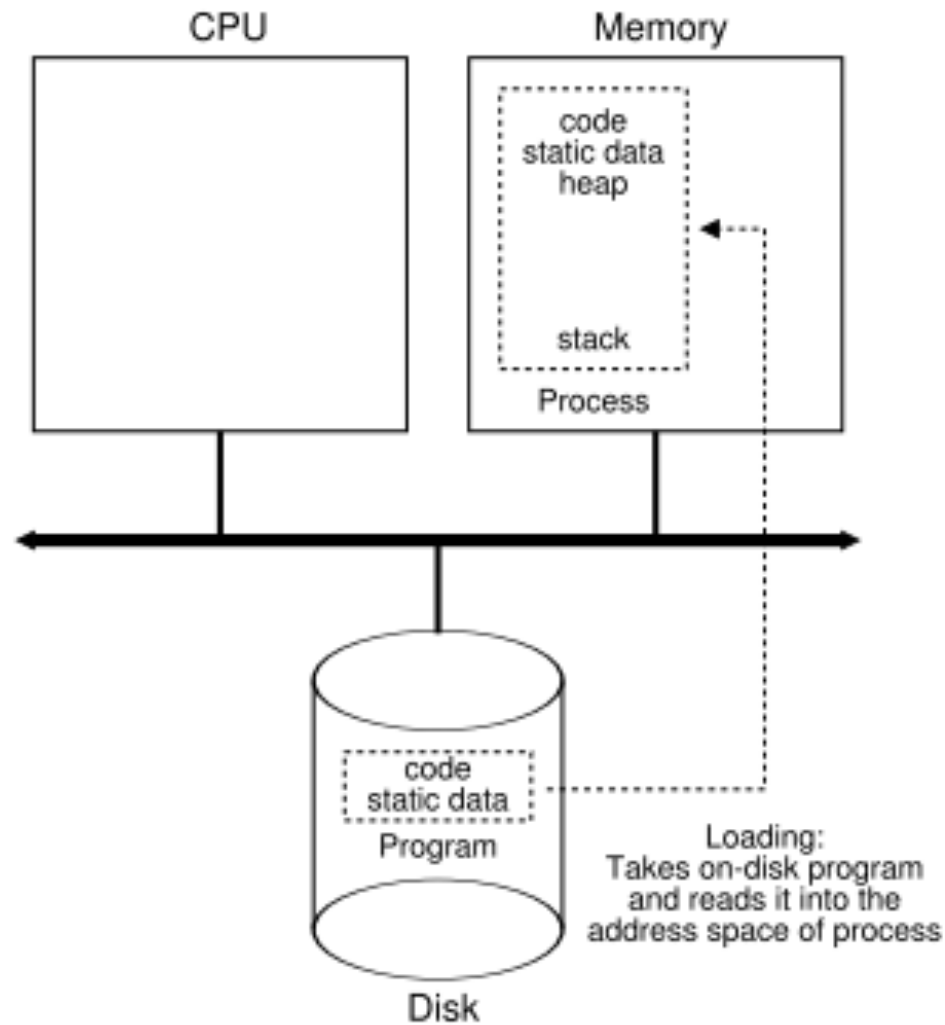


Figure 4.1: Loading: From Program To Process

# Creating a process in UNIX

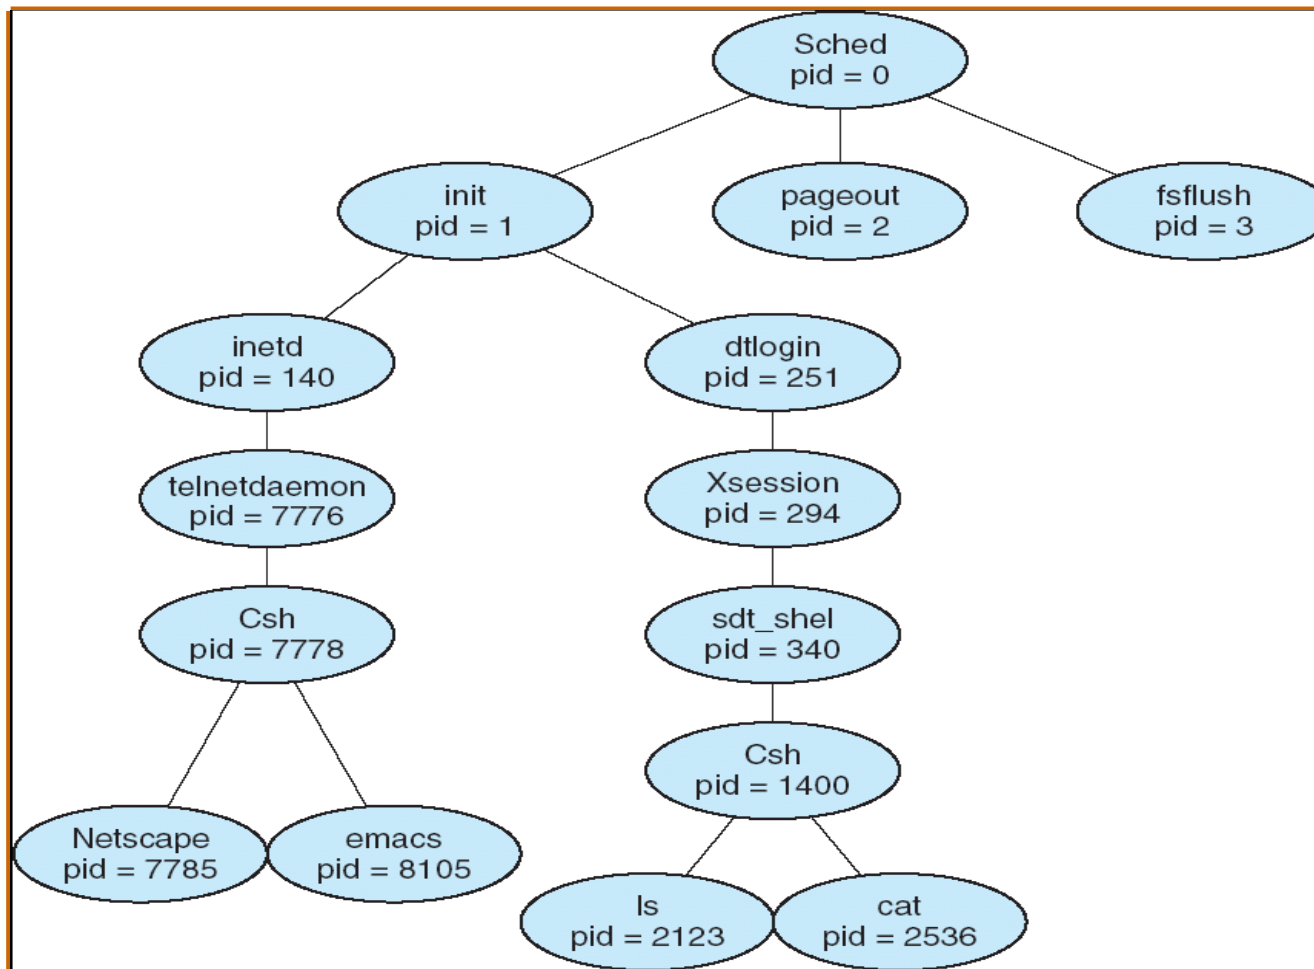
- **2 system calls: `fork()` and `exec()`**
  - When a process calls a **`fork()`** during execution, a duplicate of the calling process is created and both processes execute the next instruction after the fork.
  - When a process calls **`exec()`** with an executable as parameter, the calling process is overwritten by the process created to run this executable.
- **Do a “man” on these syscalls to find out more.**

- What really happens when you type “a.out” in the shell?
  - **The shell is itself a process.**
  - **Upon receiving this command to run a.out, the shell does a `fork()` to create a duplicate of itself.**
  - **The duplicate then does an `exec()` with the file a.out as a parameter, which results in running this program.**

OS	Program
Create entry for process list	
Allocate memory for program	
Load program into memory	
Set up stack with argc/argv	
Clear registers	
Execute <b>call</b> main()	Run main()
	Execute <b>return</b> from main
Free memory of process	
Remove from process list	

Figure 6.1: Direct Execution Protocol (Without Limits)

# A tree of processes on a typical system





# Multiple processes

- Typically, a computer system needs to manage multiple activities.
- E.g., if you are throwing a party, you not only may make a cake, but you may also bake a pizza, ....., in addition.
- Each of these activities is again a process.
- But, let us say there is no one else to help you (a single CPU).

# Options

- Perform the activities one after another (**batching**)
- Time-multiplex the CPU amongst the processes (**multiprogramming/time-sharing**), i.e., even if one activity is not fully done, you may still want to move on to processing another activity.

# Which is better?

- Say we are “batching”, i.e., finish making the cake before starting on the pizza.
- Recipe for cake: mix the flour, add sugar, place in oven, wait for 30 minutes, add icing.
- **Do you want to keep staring at the oven for 30 minutes while the cake is baking?**
- The oven is like an I/O device. Batching can result in a gross misuse of CPU resources when there is I/O.
- You want to move on to starting on the pizza while the cake is baking.

- While early mainframes employed “batching,” nearly all systems today (both desktops and servers) use time-sharing/multiprogramming.

# How do we implement time-sharing?

- We need a way of pre-empting (taking away) the CPU from the currently executing process.
- We then give the CPU to another process that can potentially use the CPU.
- This re-assigning of the CPU from one process to another is called **context-switching**.


# When should OS perform context switching?

- When an application process cannot proceed (say waiting for I/O).
- Perhaps even periodically (using timer interrupts).
- Regardless, we do NOT want applications to be aware (and require appropriate code) that they are being **context-switched**.
- **Context-switching** is transparent to an application (when you are writing code, you never care that there are other processes that may also execute in the middle).

# Implementing Context Switch

- To provide transparent context switch:
  - You need to save the “**context**” of the current process.
  - You need to restore the “**context**” of another process
- **Context** is the state of the process that is necessary for its execution
- Such saving and restoring of state ensures that the process is itself unaware that someone else executed in the middle.





# Process Control Block (PCB)

- data structure in the operating system

- Each process has a data structure called PCB
- Stores process state overall
- Including the context that is saved and restored on a context switch

pid, register value, pc, sp, hp,  
Data



# Process Data Structure

- **Code segment + Data Segment + Stack Segment + Heap Segment** – Together they are typically referred to as **address space**.
- **Process ID**
- **Parent Process**
- **Registers (context)**– note process assumes that all CPU registers are available to it.
- **Other state info maintained by OS (e.g., open files, scheduling state, etc.)**

# Process Data Structure


```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
                  RUNNABLE, RUNNING, ZOMBIE };

// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                              // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If !zero, sleeping on chan
    int killed;               // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;         // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;      // Trap frame for the
                              // current interrupt
};
```

Figure 4.5: The xv6 Proc Structure

# Context Switch



Think about what it takes  
to create a process!

# What is the advantage of a “process” view?

- Implement concurrency (between users, between activities of a user, ...)
- Insulate one activity from another

# Drawbacks of a process view

- They are heavy weight – higher scheduling (context switch) costs
  - Direct costs of switching address spaces.
  - Indirect costs (e.g., cache flushes)
- State Sharing is a problem

# What are the overheads?

- Switching code, data, stack and heap *address space*
- Saving and restoring registers
- Saving and restoring other state maintained by OS
- Can we do better?

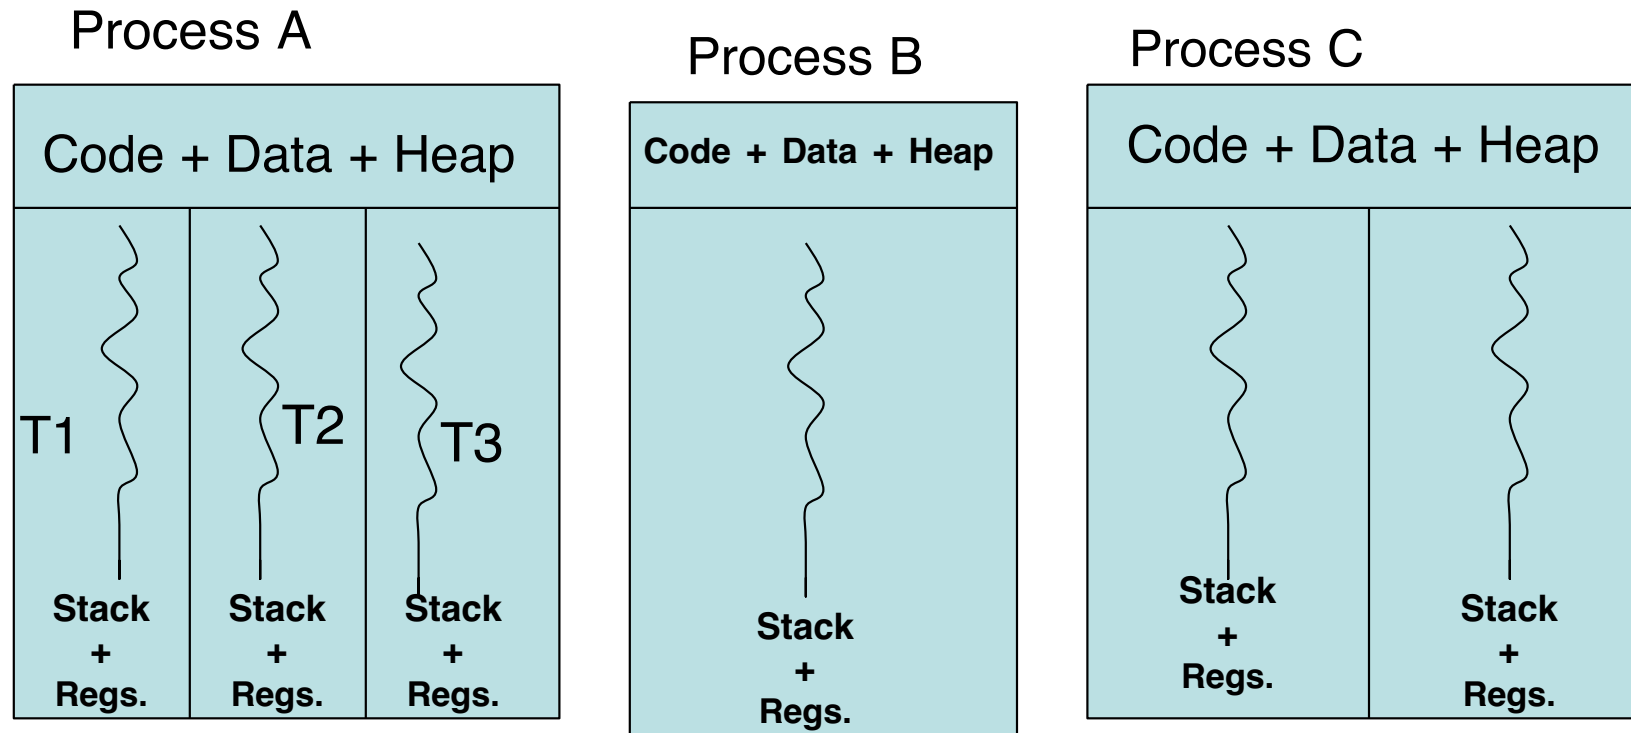
# Solution: Threads

- Activities within the same address space.
- Threads within a process share (code, data, heap).
- Only stacks are disjoint.
- Switching between threads only involves switching stacks.
- Sharing is implicit
- No protection between threads of a process – but this is OK since they are meant to be cooperative.

*stack register are not shared.*



# Threads




---

Operating System

---

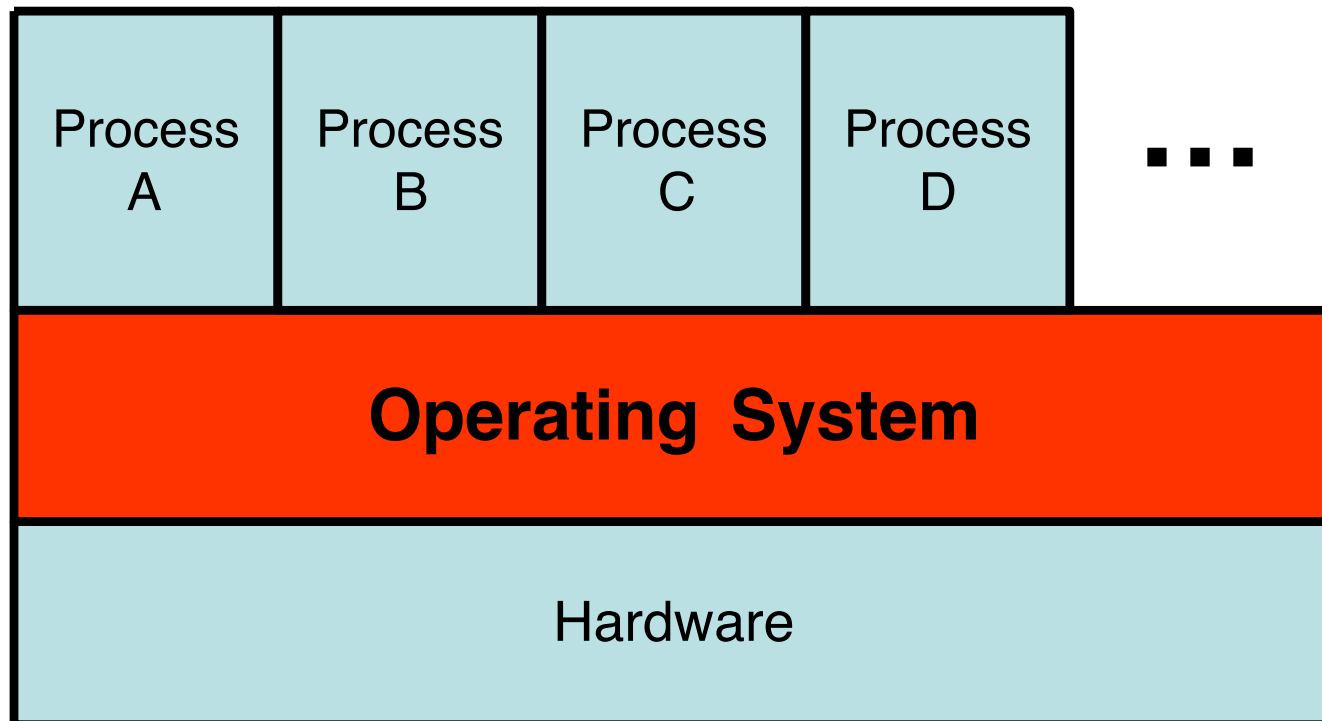
Hardware

# Solution: Threads

- More later

# Process and OS


- Focused mainly on process-to-process interactions up until now
  - but enabled by processes use of OS




# How Does the OS Help Multiple Processes Share the Same CPU?

## E.g., Two Processes on a CPU

- Let's consider (only) two processes A and B that are running on the same CPU (along with the OS)
- Let us look closely at some illuminating events in such a system



We identify four basic questions to consider



Q1: What if the process does something undesirable here?

*infinite loop  
try to reach outside of memory space  
math: divide by 0*

- What “undesirable” things might a process do?

# Undesirable #1: Executing Privileged Instructions

↘ Instructions that only execute in kernel core

- **Question:** Should a process be allowed to execute all instructions in the ISA?
- **Answer:** No
- E.g., what could go wrong if a program were allowed to execute the “halt” instruction?



# Privileged Instructions

- Instructions that are “security-sensitive” must be “privileged”
  - **Security-sensitive**: affect the operation of other process (integrity)
    - E.g., shut down computer, modify address space, modify IO
  - **Security-sensitive**: snoop data from other process (secrecy)
    - E.g., read address space, leak IO
  - **Privileged**: Run by trusted code – i.e., by the OS
  - More later...

# Undesirable #2: Certain Error Conditions

- Consider the following errors our programs often run into:
  - Segmentation fault
  - Division by zero
  - More to come

# Solution: Traps

- Let the CPU be designed s.t. upon the occurrence of the following, it enters a special error-like state and control jumps to OS
  - A process executes a “privileged” instruction
  - A process or the OS encounters one of these error conditions
- Such events are called **traps**

# Traps for system calls

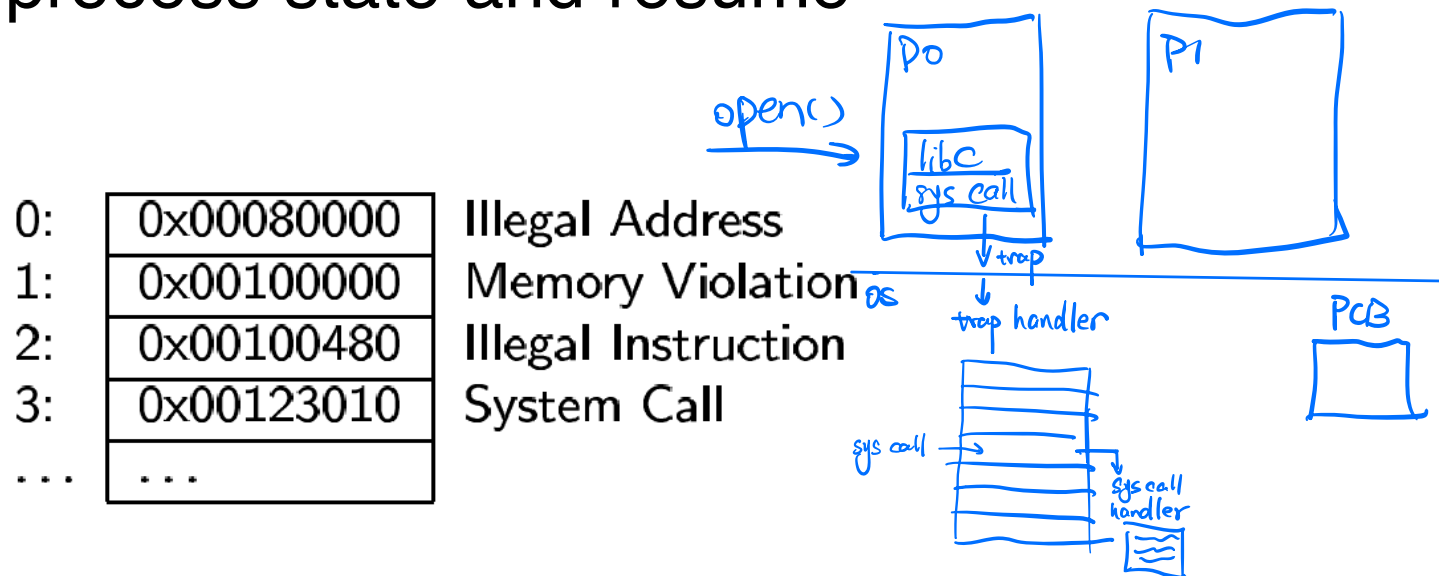
- Programs are offered a special instruction via which they can raise a trap
  - E.g., “syscall” on x86
  - Is this a privileged instruction?

?

*syscall more like unprivileged instruction, but trigger traps on purpose*

# Traps

- On detecting trap, CPU must:
  - Save process state
  - Transfer control to **trap handler** (in OS)
    - CPU indexes *trap vector* by trap number
    - Jumps to address
  - Restore process state and resume



# A Final Missing Piece!

- We would like the CPU to raise a trap when a process executes a privileged instruction
- But how would the CPU know the difference between a process and the OS?
  - An instruction is an instruction!


# Dual CPU Mode

- CPUs offer at least two “modes” of operation
  - **User mode** and **Kernel** (OS, Supervisor) mode
  - Execute privileged instruction in user mode [?] trap
  - E.g., Mode bit provided by hardware
    - Provides ability to distinguish when CPU is running process or OS
  - E.g., x86 offers four modes called “rings” with ring 0 for OS and ring 3 for processes

# Dual CPU Mode

- OS runs with CPU in kernel mode
- Is responsible to ensure programs run with CPU in user mode
- What is required to realize the above?
  - OS is the first software to run!
    - The booting up of the OS
  - OS has the ability to change CPU mode from kernel to user
  - Programs have the ability to change CPU mode from user to kernel





Q2: how does the OS get to start running here?


- Need some way outside the process's control to force control back to the OS

# Interrupts

- There must be a mechanism via which the OS gets a chance to run on the CPU every so often
  - E.g., A **timer interrupt** that periodically lets the OS run, typically, once every few milliseconds

OS

TRAP TABLE



→ 0x C75631AB

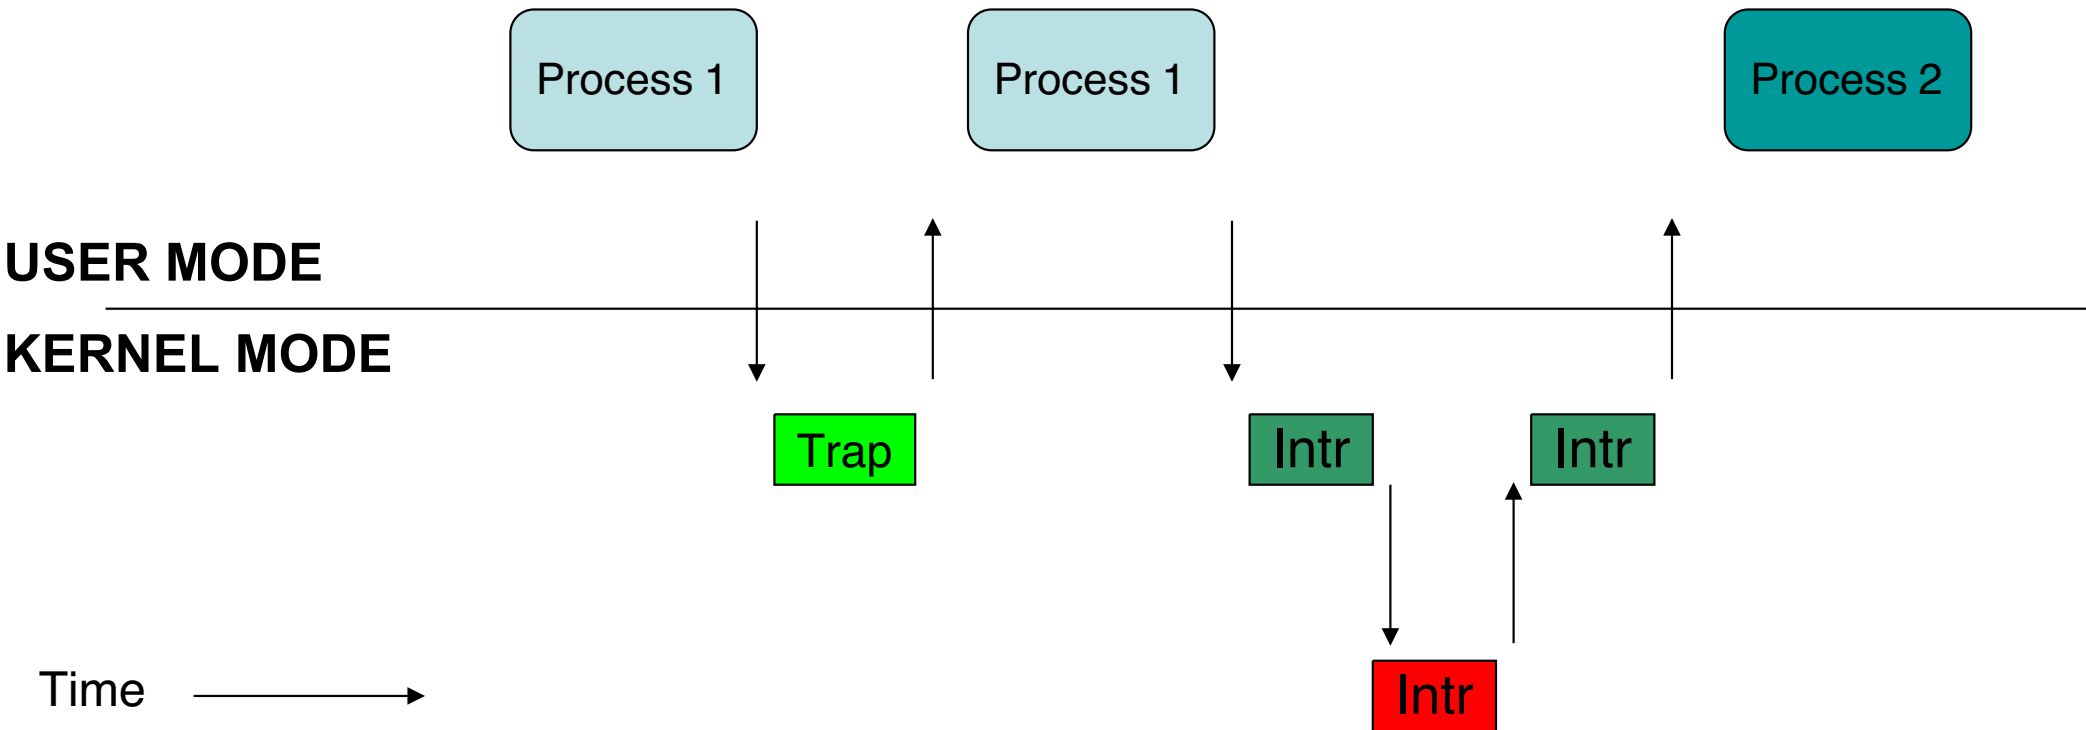
function pointer pointing to some function that will run

# Interrupts

More generally:

- Interrupts are special conditions **external to the CPU** that require OS attention
  - Note difference from traps
- CPU designed to switch <sup>from user space</sup> to kernel mode upon detecting an interrupt
  - **Example**: A keystroke raises an interrupt


# Interrupts and Traps



- Only two ways to enter supervisor mode from user mode

# Interrupts

- Are fundamental to I/O processing
  - Which we will discuss in detail later...



Q3: How do we ensure that A resumes execution at T2 as if it had not been taken off the CPU at T1?

*pick up { register state  
cell address space  
content and layout*

- By ensuring that we save the entire “state” of A at T1 and can resume it from this state at T2
- $\text{state}(A, T1) == \text{state}(A, T2)$
- What is the state of A at T1?

*knows how to access everything to all the hardware*

# State of A at time T1 (1)


- #1: Contents of A's address space
  - What are the code, data, heap, and stack values of the process at T1?

# A's Address Space


0xFFFFFFFF

virtual addresses

0x00000000



stack



heap



data



code



# State of A at time T1 (1)

- #1: Contents of A's address space
  - What are the code, data, heap, and stack values of the process at T1?
- Q: Where do these reside at time T1?
  - In a portion of main memory set aside for A
  - We rely on memory manager to ensure they remain unchanged by other processes during [T1, T2]
    - More details when we study virtual memory management

# State of A at time T1 (2)


- #2: Layout of A's address space
  - The address ranges the code, data, heap, stack span

# Layout of Address Space

0xFFFFFFFF

virtual addresses

0x00000000



# State of A at time T1 (2)

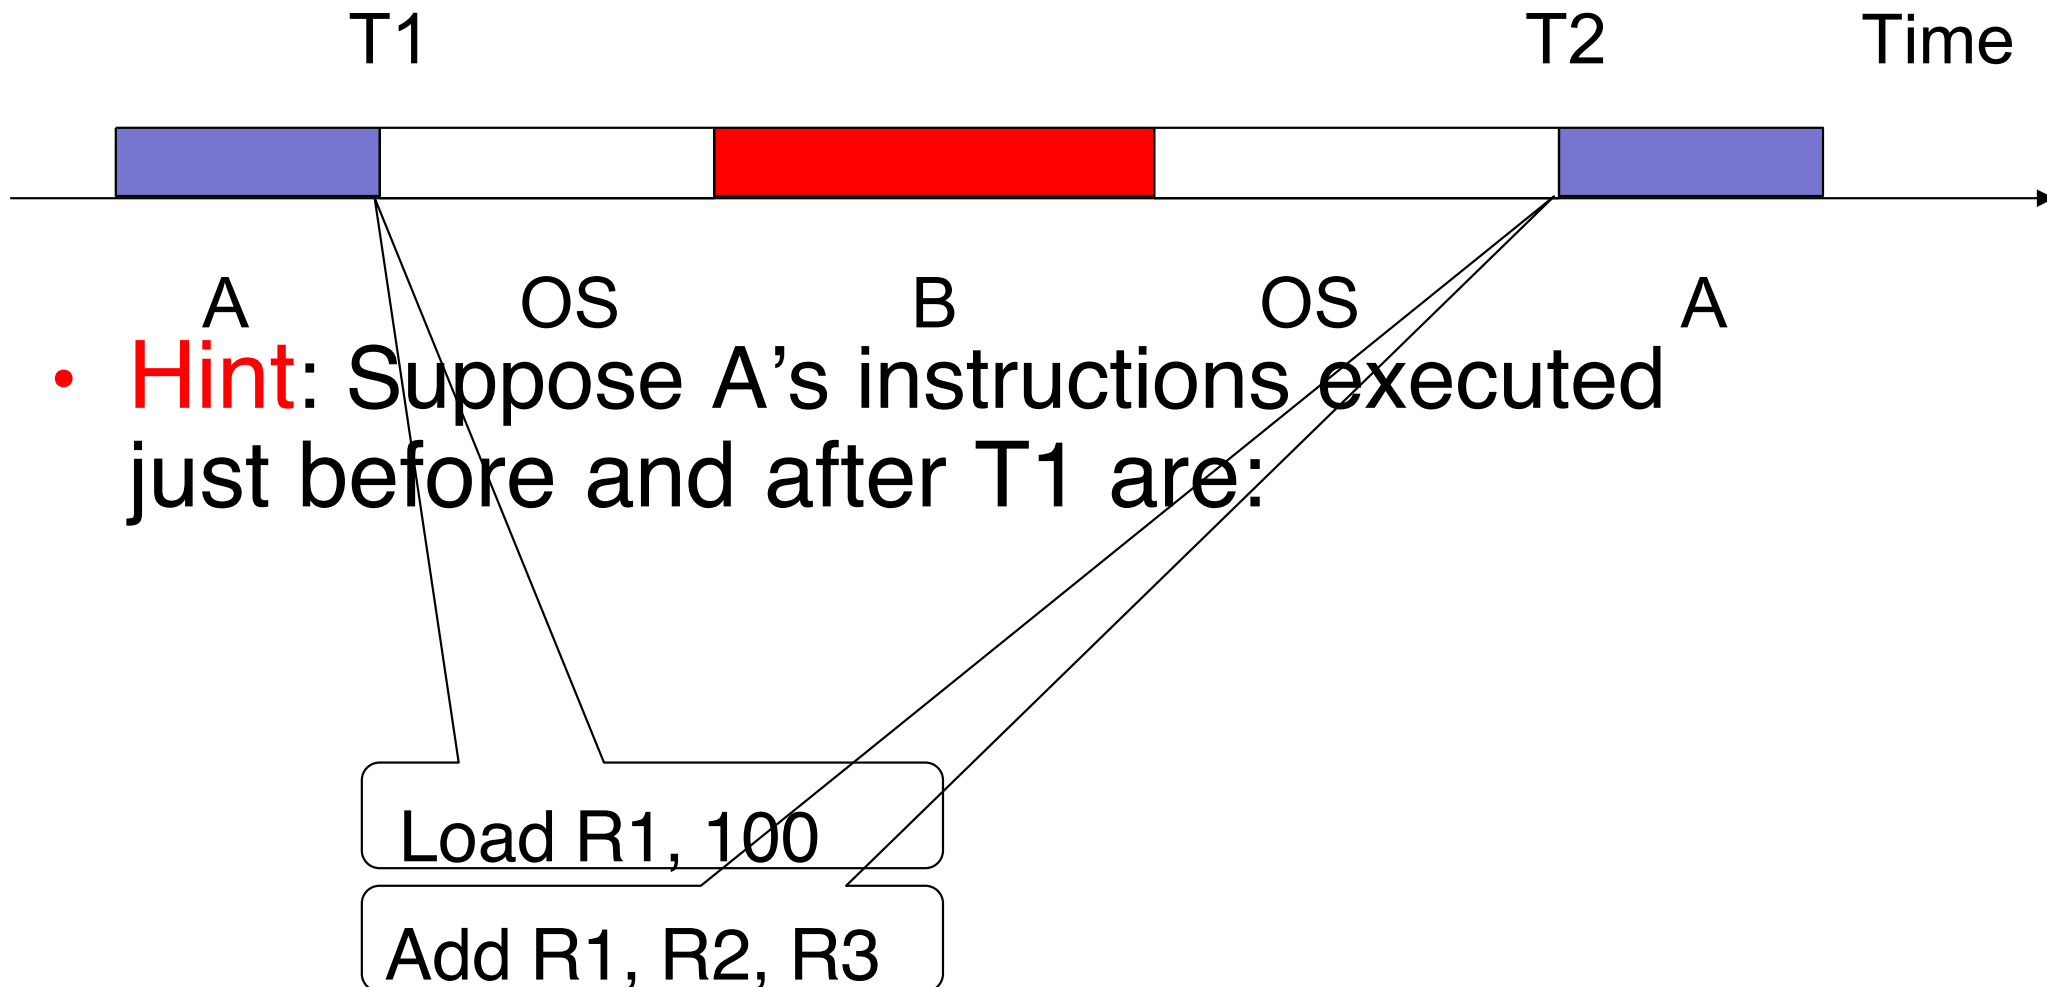
- Layout of A's address space
  - The address ranges the code, data, heap, stack span
- Q: Where are these address ranges stored?
  - Somewhere in memory
  - In whose address space? Again, A's address space is a valid choice

# State of A at time T1 (3)

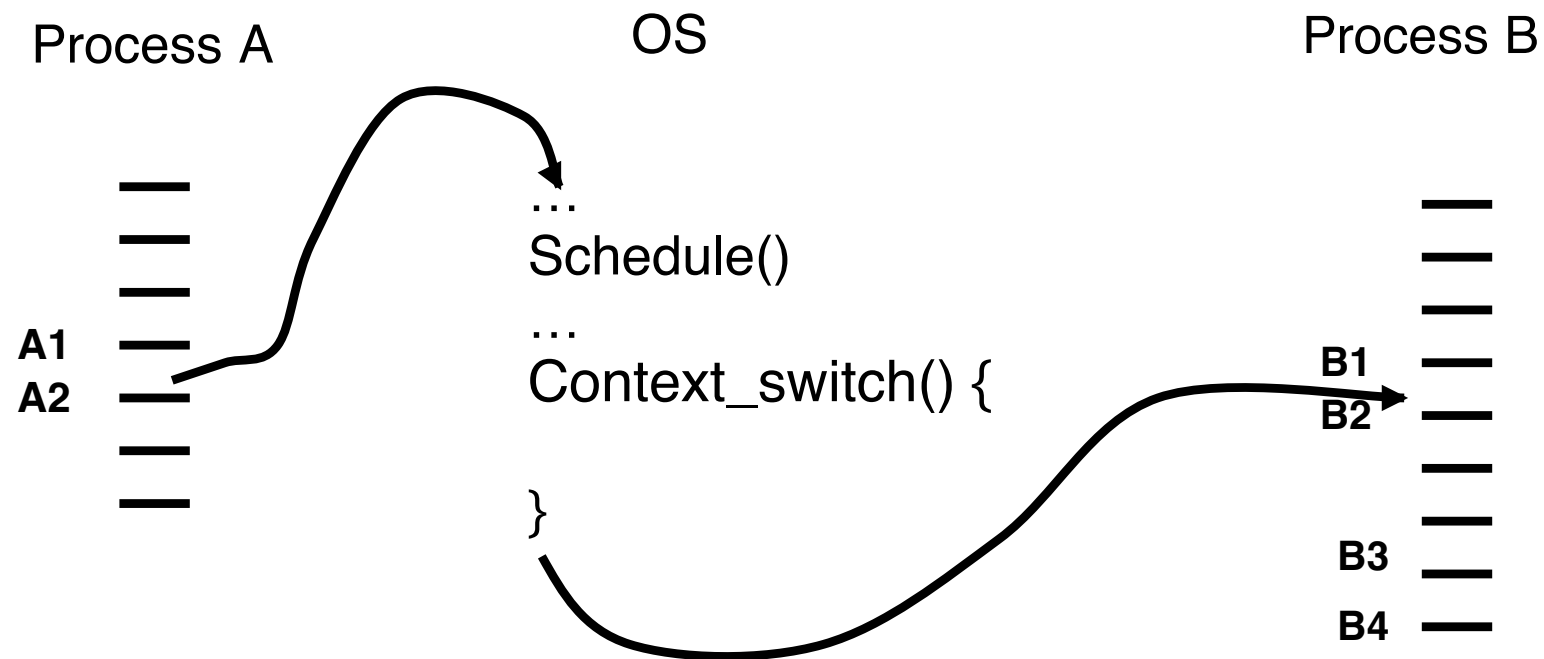
- #3: All the register values at time T1 need to be saved in main memory and restored at time T2
- Called the **hardware context** of process A
- Typically, the hardware context specifies the runtime state of the process
  - E.g., Stack Pointer Register (SP)
  - E.g., Program Counter (PC)

# State of A at time T1 (3)

- Anything else?



# Context Switch



# Context Switch: More Detail

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember addresses of... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A ...
	timer interrupt	
	save regs(A) → k-stack(A)	
	move to kernel mode	
	jump to trap handler	
Handle the trap		
Call switch() routine		
save regs(A) → proc.t(A)		
restore regs(B) ← proc.t(B)		
switch to k-stack(B)		
return-from-trap (into B)		
	restore regs(B) ← k-stack(B)	
	move to user mode	
	jump to B's PC	
		Process B ...

store in Process control block which in kernel core  
kernel → heap segment



# State of A at time T1 (4)

- #4: I/O resources being used by the process
  - E.g., open files, network sockets, etc.
- How does your process reference an open file?
  - E.g., via the *open* syscall

# State of A at time T1 (4)

- #4: I/O resources being used by the process
  - E.g., open files, network sockets, etc.
- Information held by the OS in its own address space
  - More when we discuss I/O

# Questions?