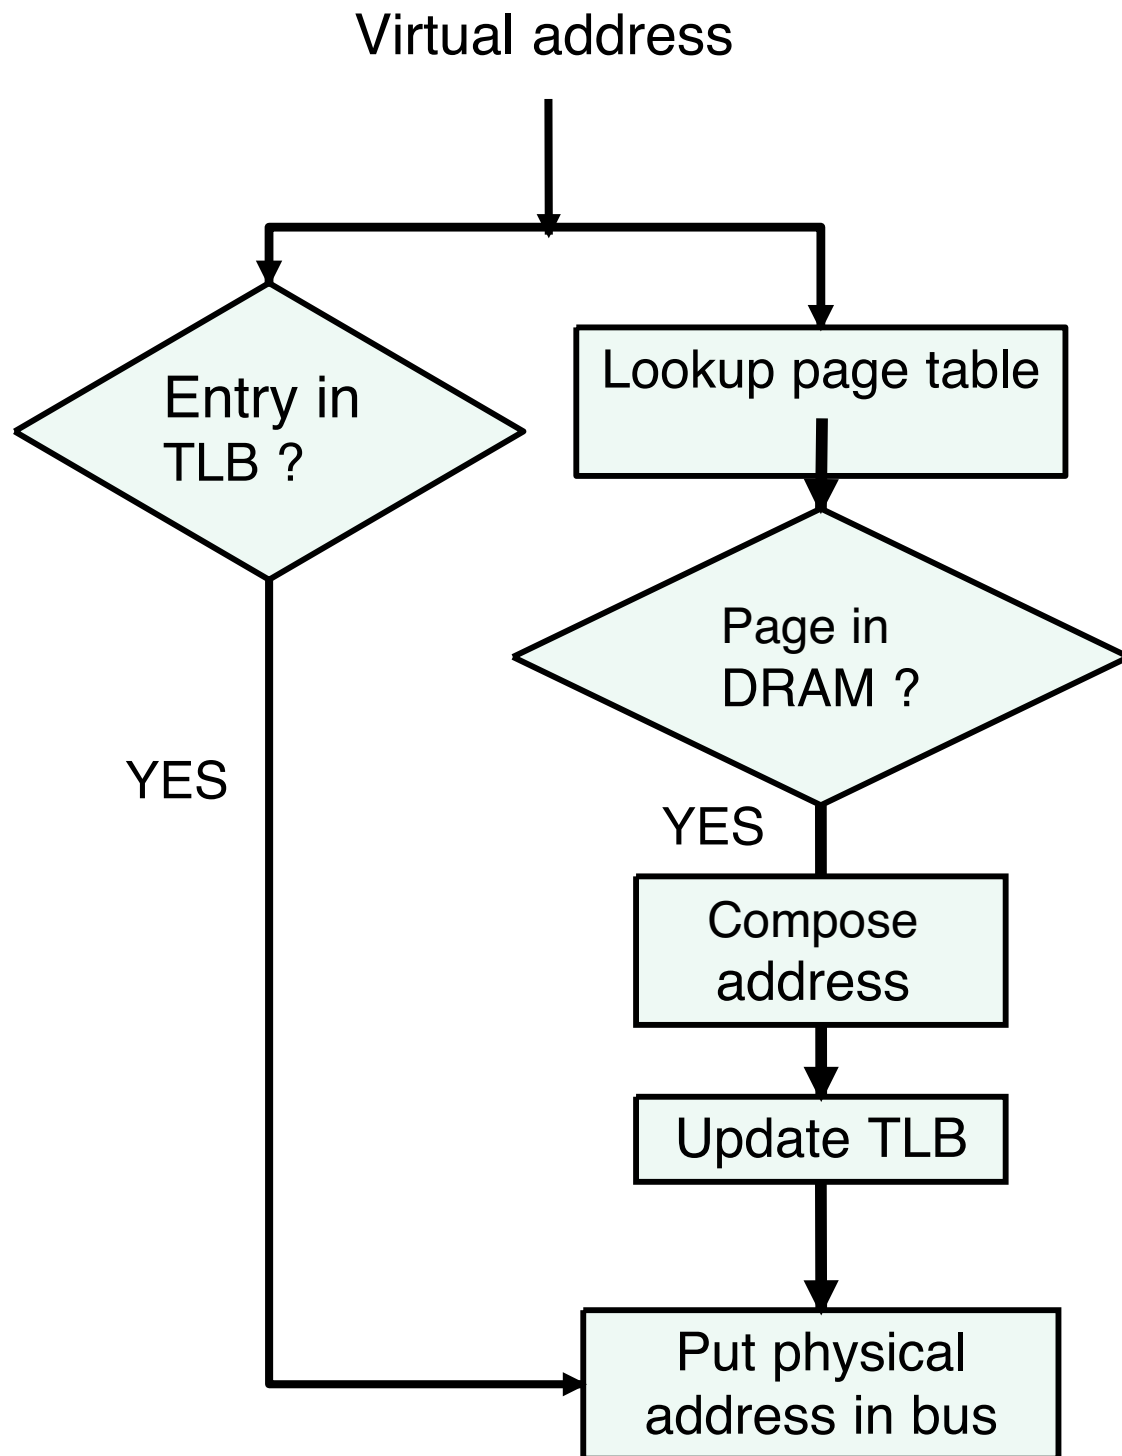



# **Virtual Memory – 2**

## **(Paging)**



# Page Fault



Interrupt Service Handler

```
PageFaultHandler() {  
    ....  
}
```

Page Fault! - Hardware Interrupt  
(H/w passes on control to S/w)

# Page-Faults

- If a Page-table mapping indicates an absence of the page in physical memory, hardware raises a “Page Fault”. — OS deal with
- OS traps this fault and the interrupt handler services the fault by initiating a disk-read request. page fault
- Once page is brought in from disk to main memory, page-table entry is updated and the process which faulted is restarted.
  - May involve replacing another page and invalidating the corresponding page-table entry.

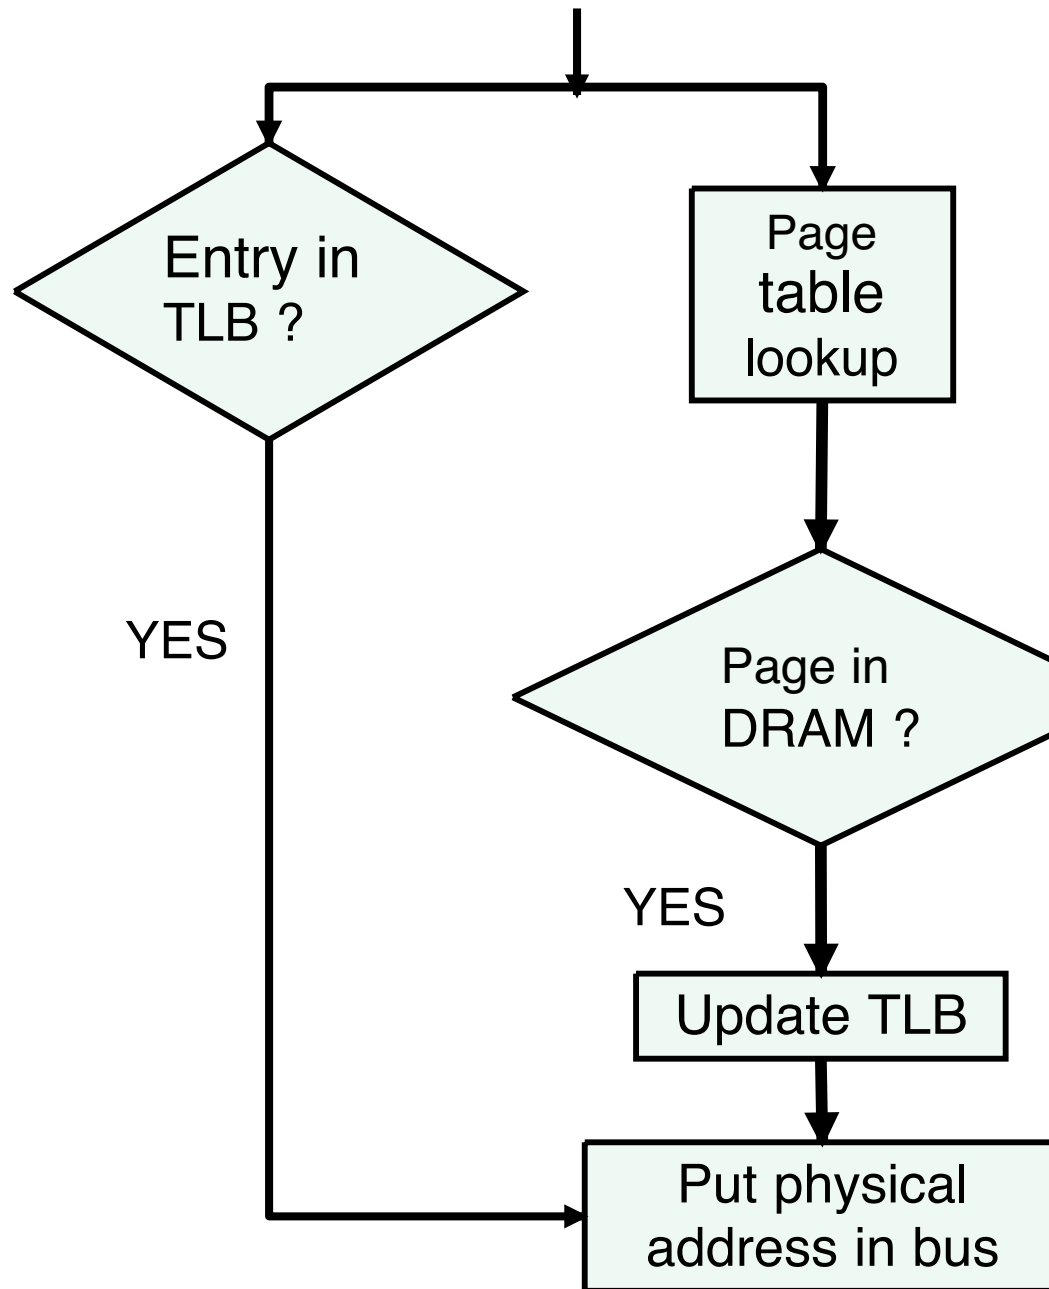


# **Physical Memory is limited!**


- **When bringing in a page, something has to be evicted.**
- **What should we evict? – page replacement algorithm.**

## Done in Hardware

Virtual address



## Done in Software



# Optimal Page Replacement Algorithm

*Find future*

- **Why optimal?**
  - No other algorithm can have # of page faults lower than this, for a given page reference stream.
- **Algorithm:**
  - At any point, amongst the given pages in memory, evict the one whose first reference from now is the furthest.

*go to future*



Optimal Page Replacement Algorithm

	3	2	1	3	4	1	6	2	4	3	4	2	1	4	5	2	1	3	4
$f_1$	3	3	3	3	4	4	4	4	4	4	4	4	4	4	5	5	5	5	5
$f_2$		2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	3	3
			1	1	1	1	6	6	6	3	3	3	1	1	1	1	1	1	4
	X	✓	X	✓	X	✓	✓	X	✓	✓	X	✓	X	✓	X	✓	✓	X	X

↓  
 page replace, look at future to find  
 the number that is furthest from now.  
 in this case, 3 is the furthest. so  
 replace 3.


# An example of OPT

**Reference String**

..... 5, 3, 3, 5, 2, 4, 4, 3, 2, .....

At this point,  
what do we replace?

**Current Physical Memory**




→ Evict

# Problem with OPT

- **Not implementable!**
- **Requires us to know the future.**
- **But it has the best page fault behavior**
- **How do we approach OPT?**

# 1. First-in First-out

- Maintain a linked list of pages in the order they were brought into physical memory.
- On a page fault, evict the one at the head.
- Put the newly brought in page (from disk) at tail of this list.
- Problems:
  - Reference String: 1,2,3,4,1,1,5,1,1,...
  - Page fault at (5) would replace (1) !
  - Need to know what is in recent use!





## 2. Least Recently Used

- Order the list of physical memory pages in decreasing order of recency of usage.
- Replace the page at the tail.
- Problem:
  - This list will need to be updated on each memory reference.
  - Asking the h/w to do this is ridiculous!
- Solution: Approximate LRU



# 3. Not Recently Used



- Referenced bit set on each Read/write by h/w
- Modified set on each write by h/w
- On startup set both R and M bits to 0.
- Periodically (using clock interrupts) the R bit is cleared. (but do not clear out modified bit).

- On a page fault, examine the state of a page
  - Class 0: R = M = 0
  - Class 1: R = 0, M = 1
  - Class 2: R = 1 M = 0
  - Class 3: R = 1 M = 1

- NRU replaces a page chosen at random from the lowest numbered nonempty class.

P\_F #1 0 0 0 0 0 0

: 1 0 0 0 0 0

PF5: 0 1 0 0 0 0

ref=1: 1 0 0 0 1 0

PF5: 1 0 0 0 1 0

Q1: Are all those steps in one pf.

Q2: what if multiple

PF2: 0 0 1 0 0 0  
 PF3: 0 0 1 0 0 0  
 PF4: 0 0 0 1 0 0

ref to one page frame  
 before page fault?

## 4. Approximate LRU using counters

when page fault, replace the  
 smallest value of counter bits.

- Keep a counter for each Phys page.
- Initially set to 0.
- At the end of each time interval (interval to be determined), shift the bits right by one position.
- Copy the reference bit to the MSB of counter and reset reference.
- For a page replacement, pick the one with the lowest counter value.
- It is an approximation of LRU because:
  - we do not differentiate between references that occurred in the same tick.
  - the history is limited by the size of the counter.


8 bit counter

start  
 0 0 0 0 0 0 0 0  
 1 0 0 0 0 0 0 0  
 0 1 0 0 0 0 0 0  
 0 0 1 0 0 0 0 0

curr 0 0 0 1 0 0 0 0  
 ref 1 0 0 0 1 0 0 0  
 curr 0 1 0 0 0 0 0 0

## 5. Second Chance Replacement or Clock Algorithm

- Same as FIFO, except you skip over the pages whose reference bit is set, resetting this bit, and moving those pages to end of list.
- Implementation:





# Summary of page replacement algorithms

- **OPT, FIFO, NRU, second-chance/clock, LRU, approximate LRU**
- **In practice, OSes use second chance/clock or some variations of it.**

## Belady's Anamoly

- **Normally you expect number of page faults to decrease as you increase physical memory size.**
- **However, this may not be the case in certain replacement algorithms**

# Example of Belady's Anamoly

- **FIFO replacement Algorithm**

- **Refernce string:**

**0 1 2 3 0 1 4 0 1 2 3 4**

- **3 physical frames**

**F F F F F F F - - F F -**

**# of faults = 9**

- **4 physical frames**

**F F F F - - F F F F F F**

**# of faults = 10**


- Algorithms which do NOT suffer from Belady's anomaly are called **stack algorithms**
- E.g. OPT, LRU.

# Modeling Paging

- **Paging behavior characterized by**
  - **Reference string**
  - **Physical memory size**
  - **Replacement algorithm**


- Visualize it as a stack (say *M*), where a page that is “referenced” is brought to the top of the stack from wherever it is.

e.g. A, B, C and D  
are virtual pages



When C is referenced ...

- Whatever is in recent use is on the top of M, and the ones that are not in recent use are at the bottom.
- In fact, the top P entries of M represent the pages in physical memory, where P is the # of physical frames.



- **Distance String:**
  - **For each element of reference string, this represents the distance of that element from the top of stack in M.**



# An example of how M changes with 5 virtual pages

## Reference String

A	B	C	D	A	B	E	A	B	C	D	E
A	B	C	D	A	B	E	A	B	C	D	E
	A	B	C	D	A	B	E	A	B	C	D
		A	B	C	D	A	B	E	A	B	C
			A	B	C	D	D	D	E	A	B
						C	C	C	D	E	A
$\infty$	$\infty$	$\infty$	$\infty$	3	3	$\infty$	2	2	4	4	4

## Distance String

## Define vector C

- **C[i] represents the number of times “i” appears in the distance string.**

## Reference String

A	B	C	D	A	B	E	A	B	C	D	E
A	B	C	D	A	B	E	A	B	C	D	E
	A	B	C	D	A	B	E	A	B	C	D
		A	B	C	D	A	B	E	A	B	C
			A	B	C	D	E	A	B	C	D
						C	C	C	D	E	A
$\infty$	$\infty$	$\infty$	$\infty$	3	3	$\infty$	2	2	4	4	4

## Distance String

C vector:  $C[0]=0$ ,  $C[1]=0$ ,  $C[2]=2$ ,  
 $C[3]=2$ ,  $C[4]=3$ ,  $C[5] \dots =0$   
 $C[\infty]=5$

# of cases

## Define Vector F

- **$F[j]$  is the number of page faults that will occur for the given reference string with “ $j$ ” physical frames.**


top =  $j$  in  $M$

- **$F[j] = C[j] + C[j+1] + C[j+2] + C[j+3] \dots + C[\infty]$**

- **It is now straightforward to prove LRU does not suffer from Belady's anomaly.**
  - **The M vector tracks what is in physical memory in the top P slots for LRU.**
  - **Note that vector C[i] is independent of physical memory size.**
  - **When you go from physical memory with j frames to (j+x) frames, note that the number of C vector terms in the RHS of equation for F decreases => Page faults can only decrease if at all!**

# Paging Issues

- Keep the essentials of what you currently need (**working set**) in physical memory.
- When something you need is not in memory, bring it in from disk:
  - On demand (demand-paging)
  - Ahead of need (pre-paging)
- Programs need to exhibit good locality to avoid “thrashing” of pages in memory.
- This usually requires good programming skills!





## Fragmentation in paging

- **Virtual Memory allows “non-contiguous” allocation of physical memory without requiring any programming changes.**
- **Note that there is only internal fragmentation of, and that too only in the last allocated page.**
- **Smaller the page, smaller the internal fragmentation.**
- **However, this reduces spatial locality.**

(sbrk)

(1) process → OS → promise  
(2) process → OS → frame

# Putting it all together!

## VAS before execution

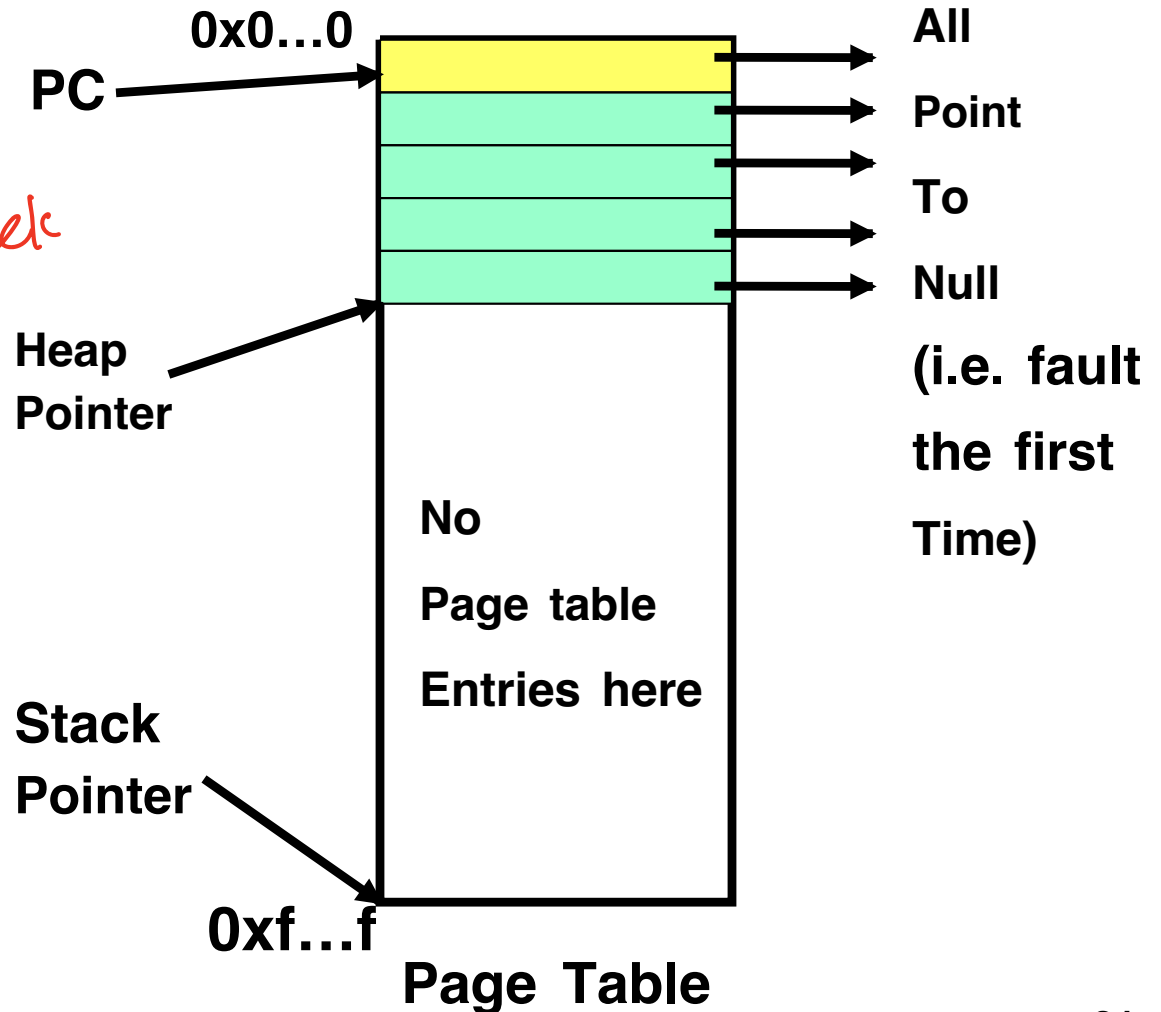
*Global - Data segment*

```
int A[8K];  
int B[8K];
```

*local - stack*

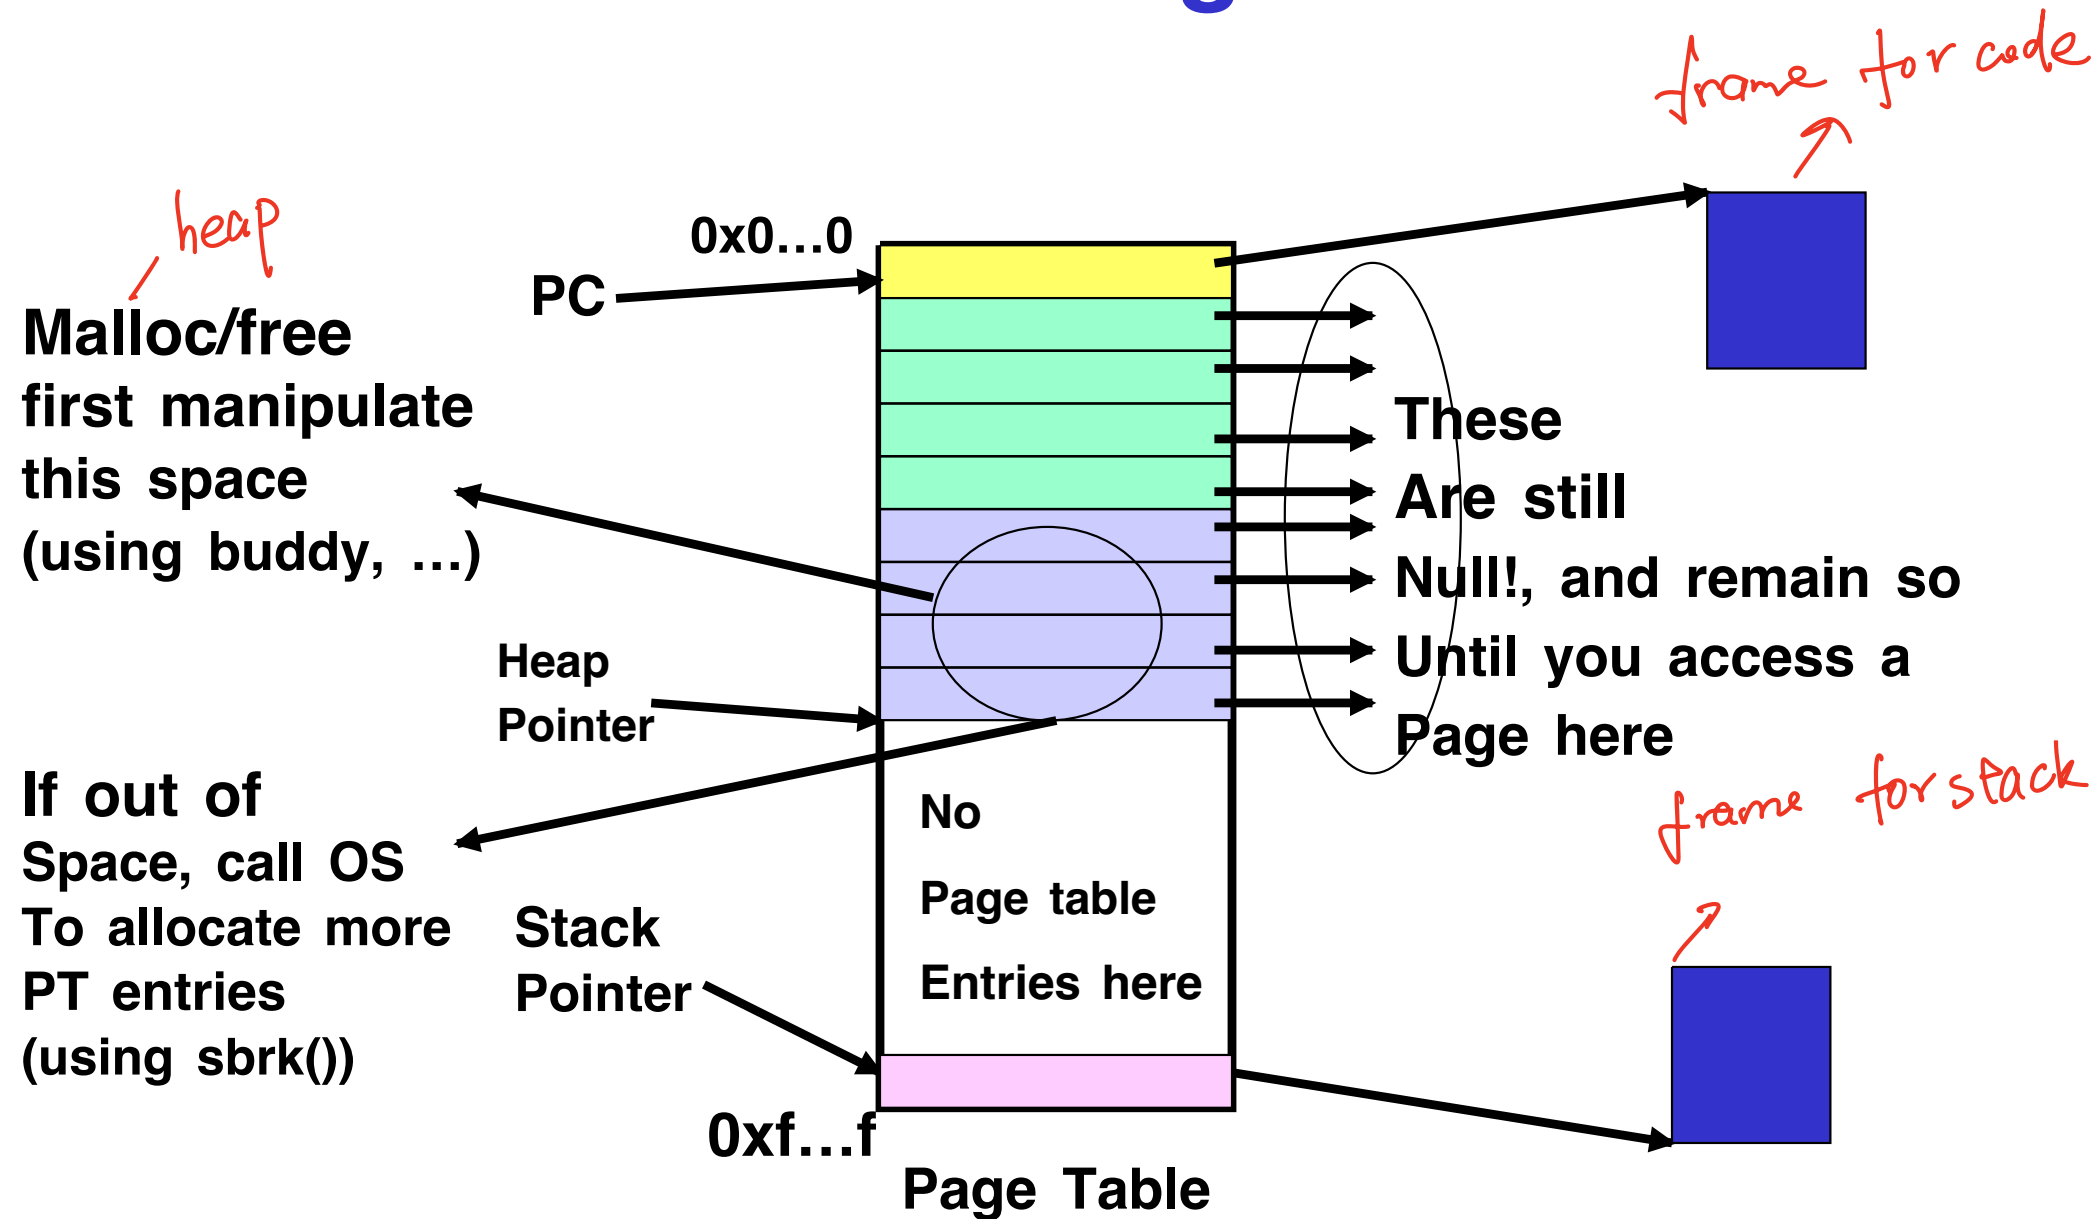
```
main() {  
    int i, j, p;  
    p = malloc(16K);  
}
```

4K page size





# After executing malloc



**Note:** you are not allocating physical memory using malloc(32

Hardware causing OS to involve is call "interrupt"  
User side causing OS to involve is call "trap"